

---

# Principles of Parallel Algorithm Design: Concurrency and Mapping

**John Mellor-Crummey**

**Department of Computer Science  
Rice University**

**`johnmc@rice.edu`**

# Last Thursday

---

- **Introduction to parallel algorithms**
  - tasks and decomposition
  - threads and mapping
  - threads versus cores
- **Decomposition techniques - part 1**
  - recursive decomposition
  - data decomposition

# Topics for Today

---

- **Decomposition techniques - part 2**
  - exploratory decomposition
  - hybrid decomposition
- **Characteristics of tasks and interactions**
- **Mapping techniques for load balancing**
  - static mappings
  - dynamic mappings
- **Methods for minimizing interaction overheads**
- **Parallel algorithm design templates**

# Exploratory Decomposition


---

- Exploration (search) of a state space of solutions
  - problem decomposition reflects shape of execution
- Examples
  - discrete optimization
    - 0/1 integer programming
  - theorem proving
  - game playing


# Exploratory Decomposition Example

## Solving a 15 puzzle


- Sequence of three moves from state (a) to final state (d)

1	2	3	4
5	6		8
9	10	7	11
13	14	15	12

(a)

1	2	3	4
5	6	7	8
9	10		11
13	14	15	12

(b)

1	2	3	4
5	6	7	8
9	10	11	
13	14	15	12

(c)

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	

(d)

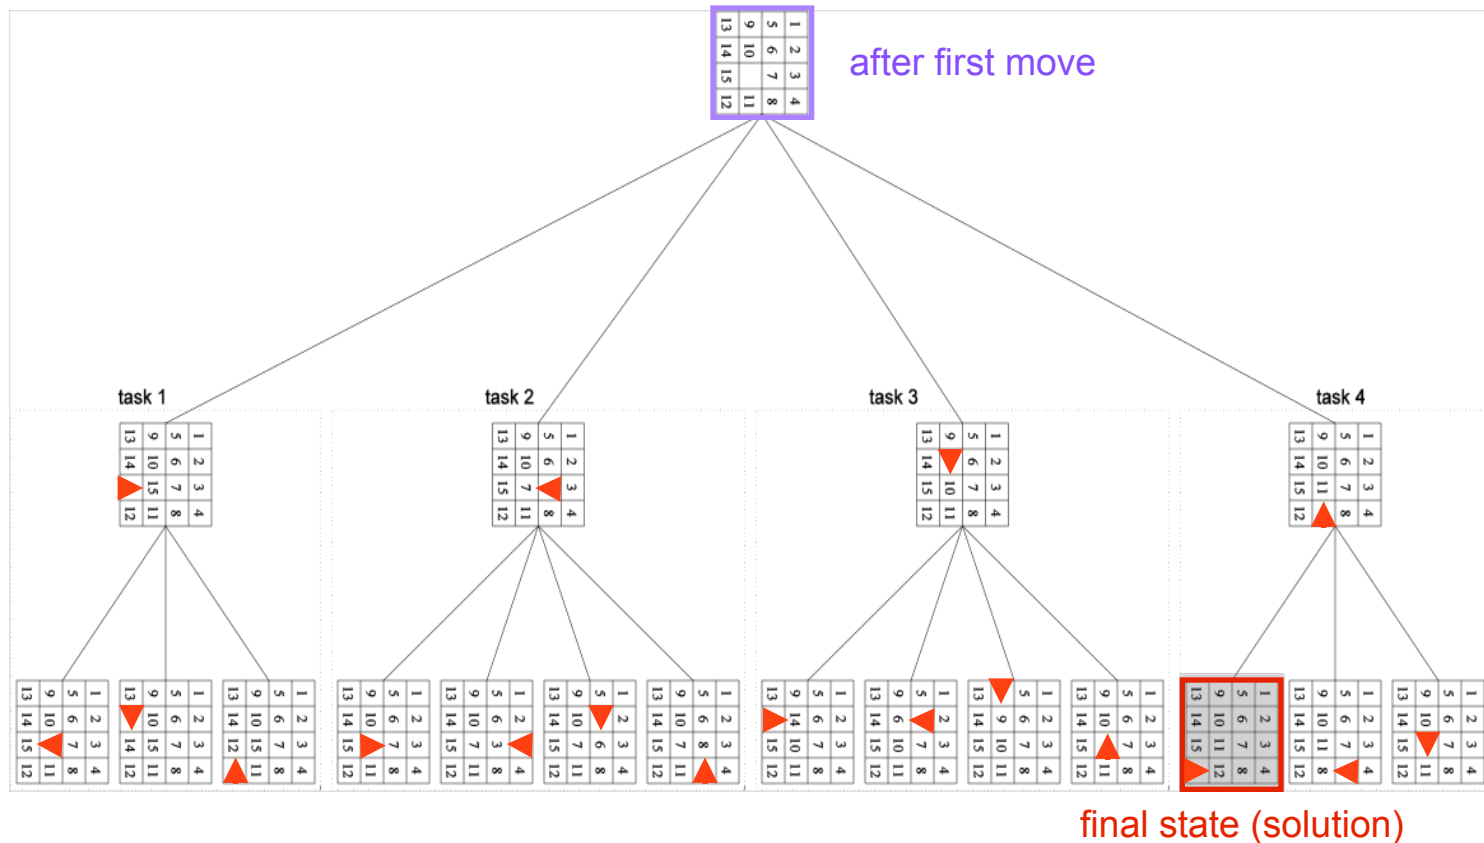
- From an arbitrary state, must search for a solution

# Exploratory Decomposition: Example

## Solving a 15 puzzle


### Search

- generate successor states of the current state
- explore each as an independent task



# Exploratory Decomposition Speedup

- Parallel formulation may perform a different amount of work



- Can cause super- or sub-linear speedup

# Speculative Decomposition

---

- Dependencies between tasks are not always known *a-priori*
  - makes it impossible to identify independent tasks
- Conservative approach
  - identify independent tasks only when no dependencies left
- Optimistic (speculative) approach
  - schedule tasks even when they may potentially be erroneous
- Drawbacks for each
  - conservative approaches
    - may yield little concurrency
  - optimistic approaches
    - may require a roll-back mechanism if a dependence is encountered



# Speculative Decomposition In Practice

---

## Discrete event simulation

- **Data structure: centralized time-ordered event list**
- **Simulation**
  - extract next event in time order
  - process the event
  - if required, insert new events into the event list
- **Optimistic event scheduling**
  - assume outcomes of all prior events
  - speculatively process next event
  - if assumption is incorrect, roll back its effects and continue

### Time Warp

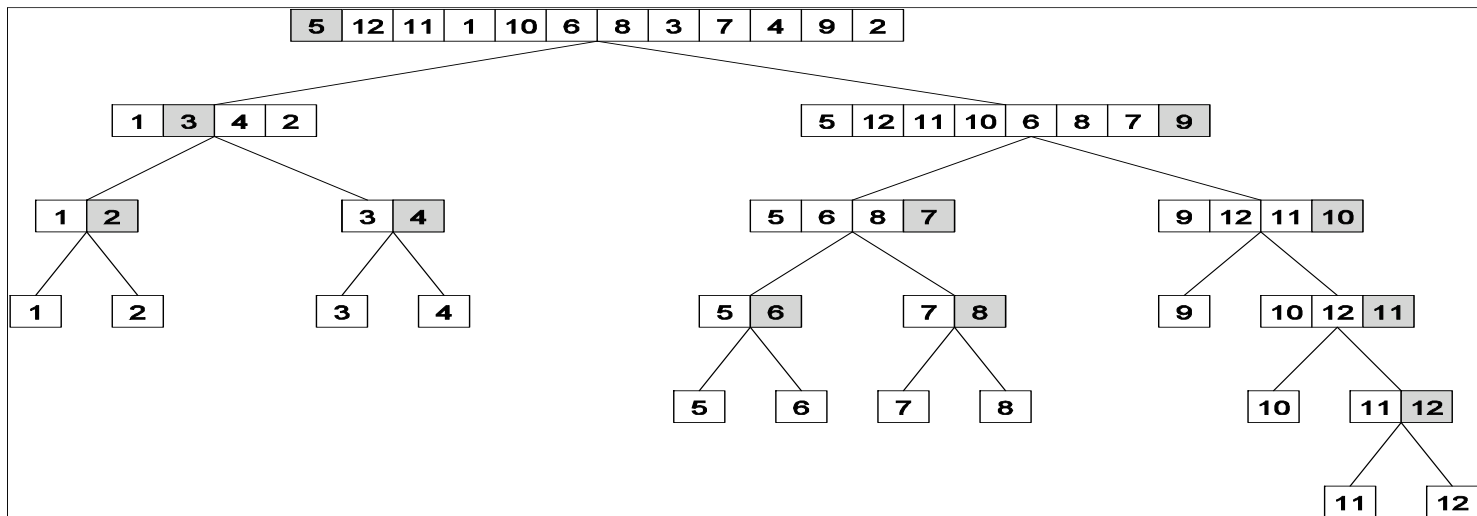
David Jefferson. "Virtual Time,"  
*ACM TOPLAS*, 7(3):404-425, July 1985

# Hybrid Decomposition

Use multiple decomposition strategies together


Often necessary for adequate concurrency

- Quicksort
  - recursive decomposition alone limits concurrency (why?)



- Climate simulation
  - data parallelism can be applied within atmosphere, ocean, land, and sea-ice simulations

# CESM Simulations on a Cray XT



Performance Limiters: Left is CAM; Right is POP.

Figure courtesy of Pat Worley (ORNL)

# Topics for Today

---

- **Decomposition techniques - part 2**
  - exploratory decomposition
  - hybrid decomposition
- ☞ **Characteristics of tasks and interactions**
  - **Mapping techniques for load balancing**
    - static mappings
    - dynamic mappings
  - **Methods for minimizing interaction overheads**
  - **Parallel algorithm design templates**

# Characteristics of Tasks

---

- **Key characteristics**
  - generation strategy
  - associated work
  - associated data size
- **Impact choice and performance of parallel algorithms**

# Task Generation

---

- **Static task generation**

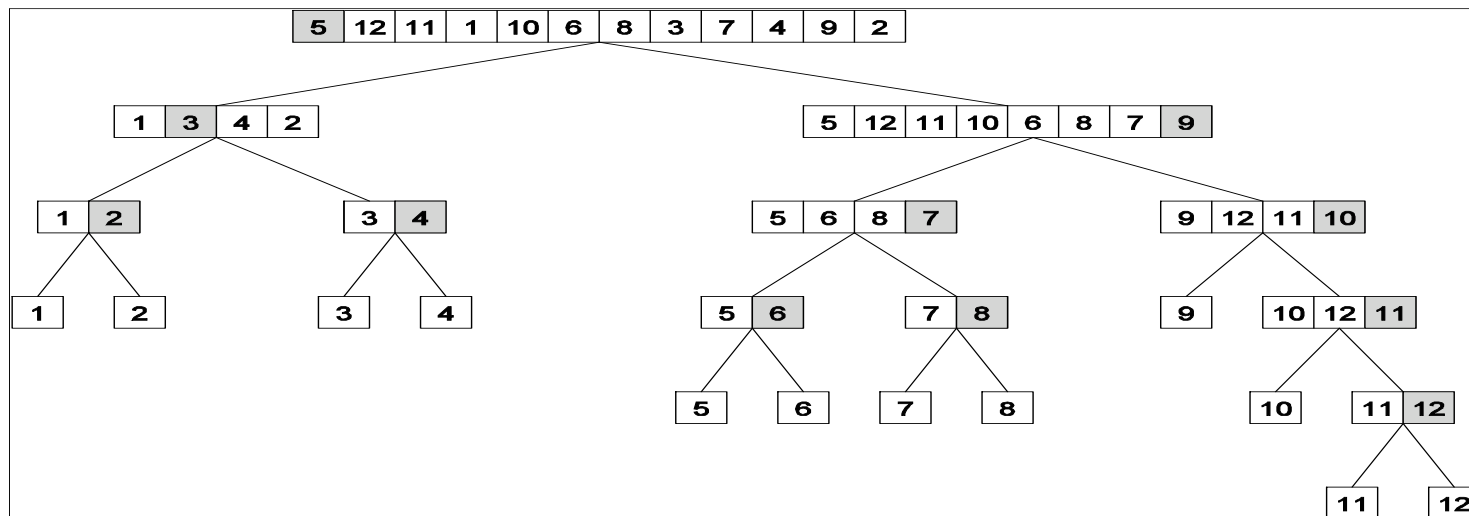
- identify concurrent tasks a-priori
- typically decompose using data or recursive decomposition
- examples
  - matrix operations
  - graph algorithms
  - image processing applications
  - other regularly structured problems

- **Dynamic task generation**

- identify concurrent tasks as a computation unfolds
- typically decompose using exploratory or speculative decompositions
- examples
  - puzzle solving
  - game playing

# Task Size

- Uniform: all the same size
  - Non-uniform
    - sometimes sizes known or can be estimated *a-priori*
    - sometimes not
      - example: tasks in quicksort
- size of each partition depends upon pivot selected



# Size of Data Associated with Tasks

---

- Data may be small or large compared to the computation
  - size(input) < size(computation), e.g., 15 puzzle
  - size(input) = size(computation) > size(output), e.g., min
  - size(input) = size(output) < size(computation), e.g., sort
- Implications
  - small data: task can easily migrate to another thread
  - large data: ties the task to a thread
    - possibly can avoid communicating the task context  
reconstruct/recompute the context elsewhere



# Characteristics of Task Interactions

---

## Orthogonal classification criteria

- **Static vs. dynamic**
- **Regular vs. irregular**
- **Read-only vs. read-write**
- **One-sided vs. two-sided**

# Characteristics of Task Interactions

---

- **Static interactions**
  - tasks and interactions are known a-priori
  - simpler to code
- **Dynamic interactions**
  - timing or interacting tasks cannot be determined a-priori
  - harder to code
    - especially using two-sided message passing APIs

# Characteristics of Task Interactions

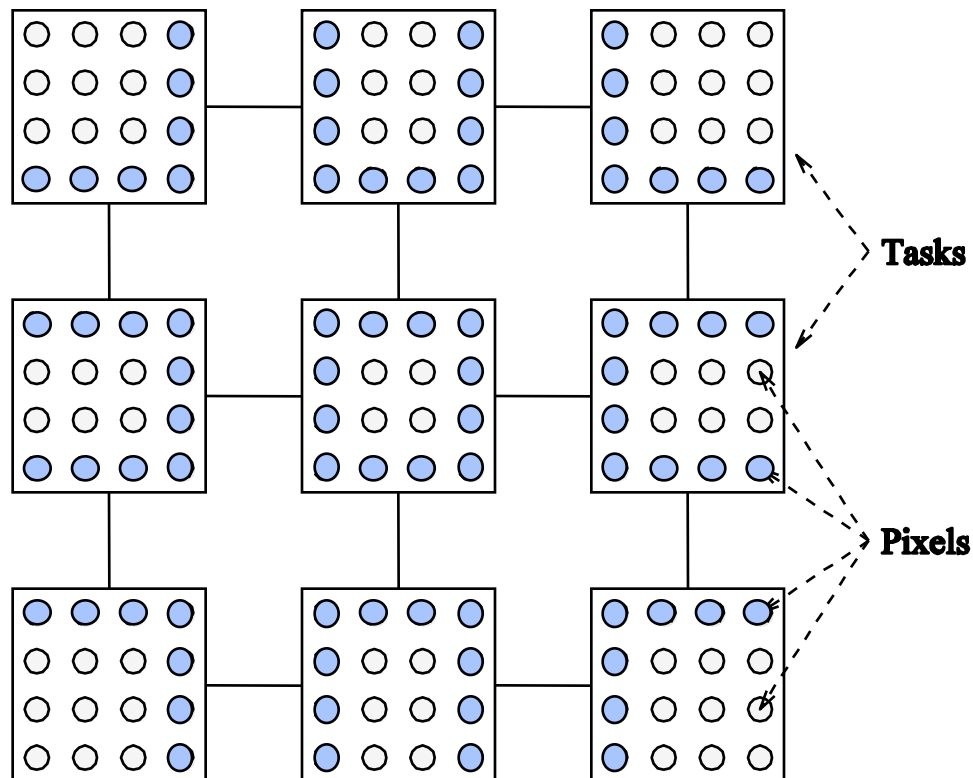
---

- **Regular interactions**
  - interactions have a pattern that can be described with a function
    - e.g. mesh, ring
  - regular patterns can be exploited for efficient implementation
    - e.g. schedule communication to avoid conflicts on network links
- **Irregular interactions**
  - lack a well-defined topology
  - modeled by a graph

# Static Regular Task Interaction Pattern

Image operations, e.g., edge detection

Nearest neighbor interactions on a 2D mesh




Sobel Edge  
Detection Stencils

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix}$$

# Static Irregular Task Interaction Pattern

**Sparse matrix-vector multiply**



# Characteristics of Task Interactions

---

- **Read-only interactions**
  - tasks only read data associated with other tasks
- **Read-write interactions**
  - read and modify data associated with other tasks
  - harder to code: requires synchronization
    - need to avoid read-write and write-write ordering races


# Characteristics of Task Interactions

---

- **One-sided**
  - initiated & completed independently by 1 of 2 interacting tasks
    - READ or WRITE
    - GET or PUT
- **Two-sided**
  - both tasks coordinate in an interaction
    - SEND and RECV

# Topics for Today

---

- **Decomposition techniques - part 2**
  - exploratory decomposition
  - hybrid decomposition
- **Characteristics of tasks and interactions**
-  **Mapping techniques for load balancing**
  - static mappings
  - dynamic mappings
- **Methods for minimizing interaction overheads**
- **Parallel algorithm design templates**



# Mapping Techniques

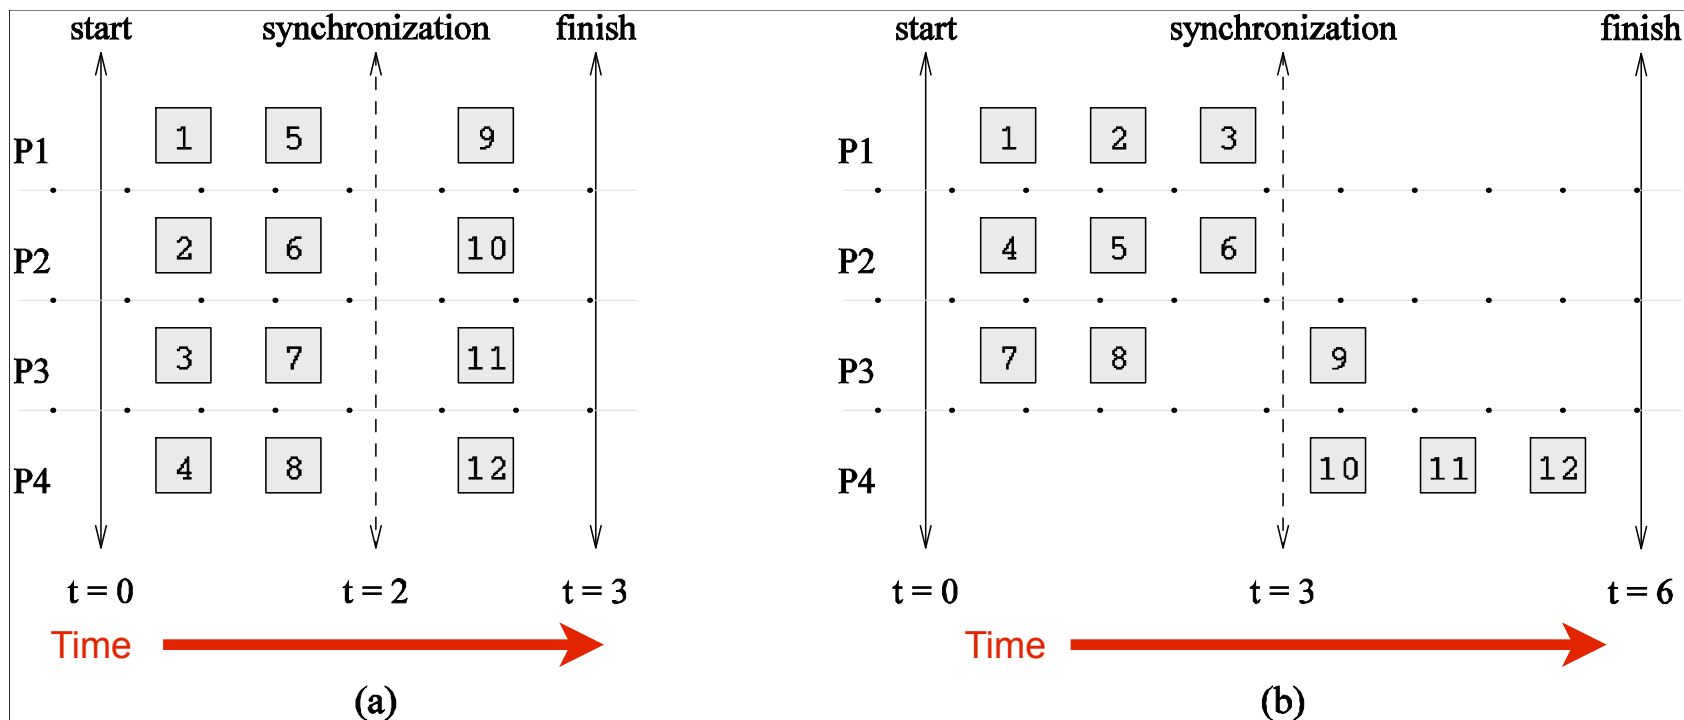
---

**Map concurrent tasks to processes for execution**

- **Overheads of mappings**
  - serialization (idling)
  - communication
- **Select mapping to minimize overheads**
- **Conflicting objectives: minimizing one increases the other**
  - assigning all work to one processor
    - minimizes communication
    - significant idling
  - minimizing serialization introduces communication

# Mapping Techniques for Minimum Idling

- Must simultaneously minimize idling and load balance
- Balancing load alone does not minimize idling



# Mapping Techniques for Minimum Idling

---

## Static vs. dynamic mappings

- **Static mapping**
  - a-priori* mapping of tasks to processes
  - requirements
    - a good estimate of task size
    - even so, optimal mapping may be NP complete
      - e.g., multiple knapsack problem
- **Dynamic mapping**
  - map tasks to processes at runtime
  - why?
    - tasks are generated at runtime, or
    - their sizes are unknown

### Factors that influence choice of mapping

- size of data associated with a task
- nature of underlying domain

# Schemes for Static Mapping

---

- **Data partitionings**
- **Task graph partitionings**
- **Hybrid strategies**

# Mappings Based on Data Partitioning


---

Partition computation using a combination of


- data partitioning
- owner-computes* rule

**Example: 1-D block distribution for dense matrices**

row-wise distribution





column-wise distribution



# Block Array Distribution Schemes

---

## Multi-dimensional block distributions



**Multi-dimensional partitioning enables larger # of processes**

# Block Array Distribution Example



---

## Multiplying two dense matrices $C = A \times B$

- Partition the output matrix  $C$  using a block decomposition
- Give each task the same number of elements of  $C$ 
  - each element of  $C$  corresponds to a dot product
  - even load balance
- Obvious choices: 1D or 2D decomposition
- Select to minimize associated communication overhead


# Data Usage in Dense Matrix Multiplication

---





# Consider: Gaussian Elimination



**Active submatrix shrinks as elimination progresses**

# Imbalance and Block Array Distributions

---

- **Consider a block distribution for Gaussian Elimination**
  - amount of computation per data item varies
  - a block decomposition would lead to significant load imbalance

# Block Cyclic Distribution

---


**Variant of the block distribution scheme that can be used to alleviate the load-imbalance and idling**

## **Steps**


- 1. partition an array into many more blocks than the number of available processes**
- 2. assign blocks to processes in a round-robin manner**
  - each process gets several non-adjacent blocks**

# Block-Cyclic Distribution

---



1D block-cyclic



2D block-cyclic

- Cyclic distribution: special case with block size = 1
- Block distribution: special case with block size is  $n/p$   
—  $n$  is the dimension of the matrix;  $p$  is the # of processes

# Decomposition by Graph Partitioning


---

## Sparse-matrix vector multiply


- Graph of the matrix is useful for decomposition
  - work  $\sim$  number of edges
  - communication for a node  $\sim$  node degree
- Goal: balance work & minimize communication
- Partition the graph
  - assign equal number of nodes to each process
  - minimize edge count of the graph partition

# Partitioning a Graph of Lake Superior

---



**Random Partitioning**



**Partitioning for minimum edge-cut**

# Mappings Based on Task Partitioning


---

## Partitioning a task-dependency graph

- Optimal partitioning for general task-dependency graph  
—NP-complete problem
- Excellent heuristics exist for structured graphs

# Mapping a Sparse Matrix

## Sparse matrix-vector product




sparse matrix structure

mapping  
partitioning



# Mapping a Sparse Matrix

## Sparse matrix-vector product



sparse matrix structure

mapping  
partitioning


17 items to  
communicate

$C0 = (4,5,6,7,8)$

$C1 = (0,1,2,3,8,9,10,11)$

$C2 = (0,4,5,6)$

13 items to  
communicate




# Hierarchical Mappings

---

- **Sometimes a single mapping is inadequate**
  - e.g., task mapping of quicksort binary tree cannot readily use a large number of processors.
- **Hierarchical approach**
  - use a task mapping at the top level
  - data partitioning within each task

# Topics for Today

---

- **Decomposition techniques - part 2**
  - exploratory decomposition
  - hybrid decomposition
- **Characteristics of tasks and interactions**
- **Mapping techniques for load balancing**
  - static mappings
  -  —dynamic mappings
- **Methods for minimizing interaction overheads**
- **Parallel algorithm design templates**

# Schemes for Dynamic Mapping

---

- **Dynamic mapping AKA dynamic load balancing**
  - load balancing is the primary motivation for dynamic mapping
- **Styles**
  - centralized
  - distributed

# Centralized Dynamic Mapping

---

- **Processes = masters or slaves**
- **General strategy**
  - when a slave runs out of work → request more from master
- **Challenge**
  - master may become bottleneck for large # of processes
- **Approach**
  - chunk scheduling: process picks up several of tasks at once
  - however
    - large chunk sizes may cause significant load imbalances
    - gradually decrease chunk size as the computation progresses

# Distributed Dynamic Mapping

---

- All processes as peers
- Each process can send or receive work from other processes
  - avoids centralized bottleneck
- Four critical design questions
  - how are sending and receiving processes paired together?
  - who initiates work transfer?
  - how much work is transferred?
  - when is a transfer triggered?
- Ideal answers can be application specific
- Cilk uses a distributed dynamic mapping: “work stealing”

# Topics for Today

---

- **Decomposition techniques - part 2**
  - exploratory decomposition
  - hybrid decomposition
- **Characteristics of tasks and interactions**
- **Mapping techniques for load balancing**
  - static mappings
  - dynamic mappings
- ☞ **Methods for minimizing interaction overheads**
  - **Parallel algorithm design templates**

# Minimizing Interaction Overheads (1)

---

## “Rules of thumb”

- **Maximize data locality**
  - don't fetch data you already have
  - restructure computation to reuse data promptly
- **Minimize volume of data exchange**
  - partition interaction graph to minimize edge crossings
- **Minimize frequency of communication**
  - try to aggregate messages where possible
- **Minimize contention and hot-spots**
  - use decentralized techniques (avoidance)



# Minimizing Interaction Overheads (2)

---

## Techniques

- **Overlap communication with computation**
  - use non-blocking communication primitives
    - overlap communication with your own computation
    - one-sided: prefetch remote data to hide latency
  - multithread code on a processor
    - overlap communication with another thread's computation
- **Replicate data or computation to reduce communication**
- **Use group communication instead of point-to-point primitives**
- **Issue multiple communications and overlap their latency**  
(reduces exposed latency)

# Topics for Today

---

- **Decomposition techniques - part 2**
  - exploratory decomposition
  - hybrid decomposition
- **Characteristics of tasks and interactions**
- **Mapping techniques for load balancing**
  - static mappings
  - dynamic mappings
- **Methods for minimizing interaction overheads**
- ☞ **Parallel algorithm design templates**

# Parallel Algorithm Model

---

- **Definition: ways of structuring a parallel algorithm**
- **Aspects of a model**
  - decomposition
  - mapping technique
  - strategy to minimize interactions

# Common Parallel Algorithm Templates

---

- **Data parallel**
  - each task performs similar operations on different data
  - typically statically map tasks to processes
- **Task graph**
  - use task dependency graph relationships to
    - promote locality, or reduce interaction costs
- **Master-slave**
  - one or more master processes generate work
  - allocate it to worker processes
  - allocation may be static or dynamic
- **Pipeline / producer-consumer**
  - pass a stream of data through a sequence of processes
  - each performs some operation on it
- **Hybrid**
  - apply multiple models hierarchically, or
  - apply multiple models in sequence to different phases

# References

---

- **Adapted from slides “Principles of Parallel Algorithm Design” by Ananth Grama**
- **Based on Chapter 3 of “Introduction to Parallel Computing” by Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar. Addison Wesley, 2003**