
Shared-memory Parallel Programming with Cilk Plus (Parts 2-3)

John Mellor-Crummey

**Department of Computer Science
Rice University**

`johnmc@rice.edu`

Last Thursday

- **Why multithreaded programming models?**
- **Threaded programming models**
- **Cilk Plus**
 - tasks
 - algorithmic complexity measures
 - scheduling
 - performance and granularity
 - task parallelism examples
 - vector addition using divide and conquer
 - nqueens: exploratory search

Outline for Today

- **Cilk Plus**
 - explore speedup and granularity
 - task parallelism example
 - **cilksort**
 - parallel loops
 - reducers
- **Data race detection with cilkscreen**
- **Assessing Cilk Plus performance with cilkview**

Review: Cilk Plus Parallel Performance Model

$$c_1 = \frac{T_1}{T_s} \quad \text{work overhead}$$

$$T_p \leq c_1 \frac{T_s}{P} + c_\infty T_\infty$$

“Minimize work overhead (c_1)
at the expense of a larger
critical path overhead (c_∞),
because work overhead
has a more direct impact
on performance”

$$T_p \approx c_1 \frac{T_s}{P} \quad \text{assuming parallel slackness}$$

Speedup Demo

- Explore speedup of naive fibonacci program

- `cp /projects/comp422/cilkplus/examples/fib ~/fib`

- `cd ~/fib`

- examine `fib.cpp` - a program for computing n^{th} fibonacci #

- build the examples: `make`

- experiment with the fibonacci program

- `make runt W=n` *computes fib(43) with n workers*

- compute `fib(43)` for different values of W , $1 \leq W \leq 8$


- what value of W yields the lowest execution time?

- what is the speedup vs. the execution time of “`./fib-serial 43`”?

- how does this compare to the total number of HW threads?

Granularity Demo

- Explore how changing increasing the granularity of parallel work in fib improves performance (by reducing c_1)
 - examine `fib-trunc.cpp` - a program for computing n^{th} fibonacci #
 - this version differs in that one can execute subtrees of height H sequentially rather than spawning parallel tasks all the way down
 - build the examples: `make`
 - experiment with the fibonacci program with truncated parallelism
 - `make runt H=h` computes `fib(43)` with lowest H levels serial
 - compute `fib(43)` for different values of H , $2 \leq H \leq 44$
 - what value of H yields the lowest execution time
 - what is the speedup vs. the execution time of “`./fib-serial 43`”?
 - how does this compare to the total number of HW threads?



Sorting in Cilk: cilksort

Variant of merge sort

- Divide array into four quarters A,B,C,D of equal size
- Sort each quarter recursively in parallel
- Merge sorted A & B into tA and C & D into tC (in parallel)
- Merge sorted tA and tC into A

High-level sketch

```
void cilksort(low,tmp,size){
    size4 = size/4
    if size <= 1 return input
    cilk_spawn cilksort(A, tA, size4);
    cilk_spawn cilksort(B, tB, size4);
    cilk_spawn cilksort(C, tC, size4);
    cilksort(D, tD, size-3*size4);
    cilk_sync;
    cilk_spawn cilkmerge(tA, tA + size4-1, tB, tB + size4-1, tA);
    cilkmerge(tC, tC + size4-1, tD, tA + size-1, tC);
    cilk_sync;
    cilkmerge(tA, tC-1, tC, tA + size-1, A);
}
```

Merging in Parallel

- How can you incorporate parallelism into a merge operation?
- Assume we are merging two sorted sequences **A** and **B**
- Without loss of generality, assume A larger than B

Algorithm Sketch

1. Find median of the elements in A and B (considered together).
2. Do binary search in A and B to find its position. Split A and B at this place to form A_1 , A_2 , B_1 , and B_2
3. In parallel, recursively merge A_1 with B_1 and A_2 with B_2

Optimizing Performance of cilksort

- Recursively subdividing all the way to singletons is expensive
- When size(remaining sequence) to sort or merge is small (2K)
 - use sequential quicksort
 - use sequential merge


Cilk Plus Parallel Loop: **cilk_for**

```
cilk_for (T v = begin; v < end; v++) {  
    statement_1;  
    statement_2;  
    ...  
}
```


- **Loop index **v****
 - type T can be an integer, ptr, or a C++ random access iterator
- **Main restrictions**
 - runtime must be able to compute total # of iterations on entry to **cilk_for**
 - must compare v with end value using <, <=, !=, >=, or >
 - loop increment must use ++, --, +=, v = v + incr, or v = v - incr
if v is not a signed integer, loop must count up
- **Implicit cilk_sync at the end of a cilk_for**

Loop with a **cilk_spawn** vs. **cilk_for**

- `for (int i = 0; i < 8; i++) { cilk_spawn work(i); } cilk_sync;`



- `cilk_for (int i = 0; i < 8; i++) { work(i);}`



Note: computation on edges

cilk_for uses
divide-and-
conquer

Restrictions for **cilk_for**

- **No early exit**
 - no break or return statement within loop
 - no goto in loop unless target is within loop body
- **Loop induction variable restrictions**
 - cilk_for (unsigned int i, j = 42; j < 1; i++, j++) { ... }**
 - only one loop variable allowed
 - cilk_for (unsigned int i = 1; i < 16; ++i) i = f();**
 - can't modify loop variable within loop
 - cilk_for (unsigned int i = 1; i < x; ++i) x = f();**
 - can't modify end within loop
 - int i; cilk_for (i = 0; i < 100; i++) { ... }**
 - loop variable must be declared in loop header

cilk_for Implementation Sketch

- Recursive bisection used to subdivide iteration space down to chunk size

```
void run_loop(first, last)
{
    if (last - first) < grainsize)
    {
        for (int i=first; i<last ++i) LOOP_BODY;
    }
    else
    {
        int mid = (last-first)/2;
        cilk_spawn run_loop(first, mid);
        run_loop(mid, last);
    }
}
```

cilk_for Grain Size

- Iterations divided into *chunks* to be executed serially
 - chunk is sequential collection of one or more iterations
- Maximum size of chunk is called *grain size*
 - grain size too small: spawn overhead reduces performance
 - grain size too large: reduces parallelism and load balance
- Default grain size
 - `#pragma cilk grainsize = min(2048, N / (8*p))`
- Can override default grain size
 - `#pragma cilk grainsize = expr`
 - `expr` is any C++ expression that yields an integral type (e.g. int, long)
e.g. `#pragma cilk grainsize = n/(4*__cilkrts_get_nworkers())`
 - pragma must immediately precede `cilk_for` to which it applies

Parallelizing Vector Addition

C

```
void vadd (real *A, real *B, int n){  
    int i; for (i=0; i<n; i++) A[i]+=B[i];  
}
```

*Cilk
Plus*

```
void vadd (real *A, real *B, int n){  
    if (n<=BASE) {  
        int i; for (i=0; i<n; i++) A[i]+=B[i];  
    } else {  
        cilk_spawn vadd (A, B, n/2);  
        vadd (A+n/2, B+n/2, n-n/2);  
        cilk_sync;  
    }  
}
```

```
void vadd (real *A, real *B, int n){  
    int i; cilk_for (i=0; i<n; i++) A[i]+=B[i];  
}
```

The Problem with Non-local Variables

- Nonlocal variables are a common programming construct
 - global variables = nonlocal variables in outermost scope
 - nonlocal = declared in a scope outside that where it is used


- Example

```
int sum = 0;
for(int i=1; i<n; i++) {
    sum += i;
}
```

- Rewriting parallel applications to avoid them is painful

Collision Detection

Automaker: hierarchical 3D CAD representation of assemblies



Computing a cutaway view

```
Node *target;
std::list<Node *> output_list;
...
void walk(Node *x) {
    switch (x->kind) {
        case Node::LEAF:
            if (target->collides_with(x))
                output_list.push_back(x);
            break;
        case Node::INTERNAL:
            for (Node::const_iterator
                child = x->begin();
                child != x->end();
                ++child)
                walk(child);
            break;
    }
}
```

Adding Cilk Plus Parallelism

Computing a cutaway view in parallel

```
Node *target;
std::list<Node *> output_list;
...
void walk(Node *x) {
    switch (x->kind) {
    case Node::LEAF:
        if (target->collides_with(x))
            output_list.push_back(x);
        break;
    case Node::INTERNAL:
        cilk_for (Node::const_iterator
                 child = x->begin();
                 child != x->end();
                 ++child)
            walk(child);
        break;
    }
}
```

**Global variable
causes data races!**

Solution 1: Locking

Computing a cutaway view in parallel

```
Node *target;
std::list<Node *> output_list;
mutex m;
...
void walk(Node *x) {
    switch (x->kind) {
    case Node::LEAF:
        if (target->collides_with(x))
            { m.lock(); output_list.push_back(x); m.unlock(); }
        break;
    case Node::INTERNAL:
        cilk_for (Node::const_iterator
                    child = x->begin();
                    child != x->end();
                    ++child)
            walk(child);
        break;
    }
}
```

- Add a mutex to coordinate accesses to output_list
- Drawback: lock contention can hurt parallelism

Solution 2: Refactor the Code

```
Node *target;
std::list<Node *> output_list;
...
void walk(Node *x, std::list<Node *> &o_list) {
    switch (x->kind) {
    case Node::LEAF:
        if (target->collides_with(x))
            o_list.push_back(x);
        break;
    case Node::INTERNAL:
        std::vector<std::list<Node *>>
            child_list(x.num_children);
        cilk_for (Node::const_iterator
            child = x->begin();
            child != x->end();
            ++child)
            walk(child, child_list[child]);
        for (int i=0; i < x.num_children; ++i)
            o_list.splice(o_list.end(), child_list[i]);
        break;
    }
}
```

- **Have each child accumulate results in a separate list**
- **Splice them all together**
- **Drawback: development time, debugging**

Solution 3: Cilk Plus Reducers

```
Node *target;

cilk::reducer_list_append<Node *> output_list;
...
void walk(Node *x) {
    switch (x->kind) {
        case Node::LEAF:
            if (target->collides_with(x))
                output_list.push_back(x);
            break;
        case Node::INTERNAL:
            cilk_for (Node::const_iterator
                    child = x->begin();
                    child != x->end();
                    ++child)
                walk(child);
            break;
    }
}
```


- **Resolve data races without locking or refactoring**
- **Parallel strands may see different views of reducer, but these views are combined into a single consistent view**

Cilk Plus Reducers


- Reducers support update of nonlocal variables without races
 - deterministic update using associative operations
 - e.g., global sum, list and output stream append, ...
 - result using is same as serial version
 - independent of # processors or scheduling
- Can be used without significant code restructuring
- Can be used independently of the program's control structure
 - unlike constructs defined only over loops
- Implemented efficiently with minimal overhead
 - they don't use locks in their implementation
 - avoids loss of parallelism from enforcing mutual exclusion when updating shared variables

Reducers

Serial execution (depth first):



Parallel execution:




Matteo Frigo, Pablo Halpern, Charles E. Leiserson, Stephen Lewin-Berlin, Reducers and other Cilk++ hyperobjects. Slides for SPAA'09, August 11–13, 2009, Calgary, Alberta, Canada.

Reducing Over List Concatenation

Program:

```
x.append(0);  
cilk_spawn x.append(1);  
x.append(2);  
x.append(3);  
cilk_sync;
```

Serial execution:




Matteo Frigo, Pablo Halpern, Charles E. Leiserson, Stephen Lewin-Berlin, Reducers and other Cilk++ hyperobjects. Slides for SPAA'09, August 11–13, 2009, Calgary, Alberta, Canada.

Reducing Over List Concatenation

Program:

```
x.append(0);  
cilk_spawn x.append(1);  
x.append(2);  
x.append(3);  
cilk_sync;
```

Parallel execution:



Matteo Frigo, Pablo Halpern, Charles E. Leiserson, Stephen Lewin-Berlin, Reducers and other Cilk++ hyperobjects. Slides for SPAA'09, August 11–13, 2009, Calgary, Alberta, Canada.

Using Cilk Plus Reducers

- Include the appropriate Cilk Plus reducer header file

`reducer_opadd.h, reducer_min.h, reducer_max.h,
reducer_opor.h, reducer_opand.h, reducer_opxor,
reducer_list.h, reducer_ostream.h`

- Declare a variable as a reducer rather than a standard type

—global sum

– `cilk::reducer_opadd<unsigned long> sum`

—list reducer

– instead of “`std::list<int> sequence`”, use
`cilk::reducer_list_append<int> sequence`

- Use reducers in the midst of work that includes parallelism created with `cilk_spawn` or `cilk_for`
- Retrieve the reducer's terminal value with `var.get_value()` after the parallel updates to the reducer are complete

reducer Demo

- Explore benefits of reducers

- cp /projects/comp422/cilkplus/examples/sum ~/sum

- cd ~/sum

- inspect the programs

- race.cpp - a cilk_for summation with a race

- lock.cpp - a safe cilk_for summation using a lock

- reducer.cpp - a safe cilk_for summation using a reducer

- make run

- On STIC,
see /projects/comp422/cilkplus/examples/features-tutorial
 - In that directory, there is subdirectory “reducers” that
contains more reducer examples

Cilk++ Reducers Under the Hood

- If no steal occurs, a reducer behaves like a normal variable
- If a steal occurs
 - the continuation receives a view with an identity value
 - the child receives the reducer as it was prior to the spawn
 - at the corresponding `cilk_sync`
 - the value in the continuation is merged into the reducer held by the child using the reducer's `reduce` operation
 - the new view is destroyed
 - the original (updated) object survives

Concurrency Cautions

- Only limited guarantees between descendants or ancestors
 - DAG precedence order maintained and nothing more
 - don't assume atomicity between different procedures!

Race Conditions

- **Data race**
 - two parallel strands access the same data
 - at least one access is a write
 - no locks held in common
- **General determinacy race**
 - two parallel strands access the same data
 - at least one access is a write
 - a common lock protects both accesses

A Data Race Example

- Example

```
int sum = 0;  
cilk_for(int i=1; i<n; i++) {  
    sum += i;  
}
```

- What can go wrong?

—concurrent reads and writes can interleave in unpredictable ways

time ↓

	read sum
	read sum
	write sum + i_j
	write sum + i_k

legend
thread n
thread m

— the update by thread m is lost!

Cilkscreen

- Detects and reports data races when program terminates
 - finds all data races even those by third-party or system libraries
- Does not report determinacy races
 - e.g. two concurrent strands use a lock to access a queue
 - enqueue & dequeue operations could occur in different order
 - potentially leads to different result

Race Detection Strategies in Cilkscreen

- **Lock covers**
 - two conflicting accesses to a variable don't race if some lock L is held while each of the accesses is performed by a strand
- **Access precedence**
 - two conflicting accesses do not race if one must precede the other
 - access A is by a strand X, which precedes the `cilk_spawn` of strand Y which performs access B
 - access A is performed by strand X, which precedes a `cilk_sync` that is an ancestor of strand Y

Cilkscreen Race Example

```
#include <stdio.h>
#include "mutex.h"

long sum = 0;
mutex m;

#ifdef SYNCH
#define LOCK m.lock()
#define UNLOCK m.unlock()
#else
#define LOCK
#define UNLOCK
#endif

void do_accum(int l, int u)
{
    if (u == l) { LOCK; sum += l; UNLOCK; }
    else {
        int mid = (u+l)/2;
        cilk_spawn do_accum(l, mid);
        do_accum(mid+1, u);
    }
}

int main()
{
    do_accum(0, 1000);
    printf("sum = %d\n", sum);

    long ssum = 0;
    for (int i = 0; i <= 1000; i++) ssum +=i;
    printf("serial sum = %d\n", ssum);
}
```

note: mutex class coded using pthread_mutex lock primitives

Cilkscreen Limitations

- Only detects races between Cilk Plus strands
 - depends upon their strict fork/join paradigm
- Only detects races that occur given the input provided
 - does not prove the absence of races for other inputs
 - choose your testing inputs carefully!
- Runs serially, 15-30x slower
- Increases the memory footprint of an application
 - could cause an error if memory demand is too large
- If you build your program with debug information (compile with -g), cilkscreen will associate races with source line numbers

Cilkscreen Output

Cilkscreen Race Detector V2.0.0, Build 3229
summing integers from 0 to 20000

Race condition on location 0x6016f0

write access at 0x400b7f: (/home/johnmc/examples/races/sum2.c:22, do_accum+0x169)

read access at 0x400b78: (/home/johnmc/examples/races/sum2.c:22, do_accum+0x162)

called by 0x400ca9: (/home/johnmc/examples/races/sum2.c:26, do_accum+0x293)

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

...

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

called by 0x400e47: (/home/johnmc/examples/races/sum2.c:37, main+0x85)

Race condition on location 0x6016f0

write access at 0x400b7f: (/home/johnmc/examples/races/sum2.c:22, do_accum+0x169)

write access at 0x400b7f: (/home/johnmc/examples/races/sum2.c:22, do_accum+0x169)

called by 0x400ca9: (/home/johnmc/examples/races/sum2.c:26, do_accum+0x293)

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

...

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

called by 0x400c8f: (/home/johnmc/examples/races/sum2.c:25, do_accum+0x279)

called by 0x400e47: (/home/johnmc/examples/races/sum2.c:37, main+0x85)

sum = 200010000

serial sum = 200010000

2 errors found by Cilkscreen

Cilkscreen suppressed 119998 duplicate error messages

cilkscreen Demo

- Explore cilkscreen race detection

- cp /projects/comp422/cilkplus/examples/races ~/races

- cd ~/races

- inspect the programs

- sum1.c - a cilk_for summation with a race

- sum2.c - a task parallel summation w/ optional mutex

- nocover.c - program using locks with races

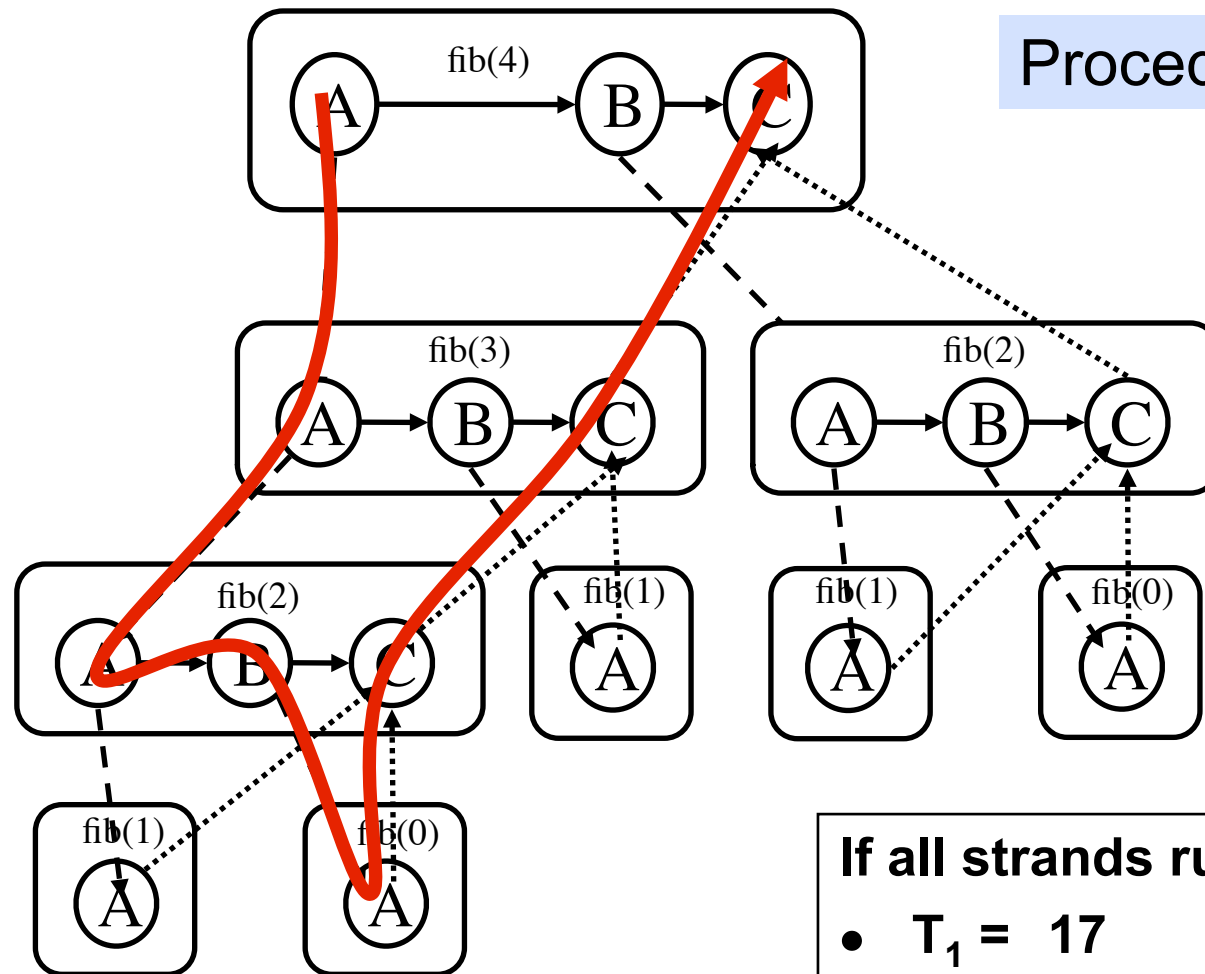
- cover.c - lock cover prevents races

- make run

Performance Measures

- T_s = serial execution time
- T_1 = execution time on 1 processor (total work), $T_1 \geq T_s$
- T_p = execution time on P processors
- T_∞ = execution time on infinite number of processors
 - longest path in DAG
 - length reflects the cost of computation at nodes along the path
 - known as “critical path length”

Work and Critical Path Example




Procedure oriented view

If all strands run in unit time

- $T_1 = 17$
- $T_\infty = 8$ (critical path length)

Execution DAG View


- Cilk Plus uses the word “strand” for a serial section of the program
- A “knot” is a point where three or more strands meet
- Two kinds of knots
 - spawn knots: one input strand, two output strands
 - sync knots: two or more input strands, one output strand



```
...  
do_stuff1();  
cilk_spawn func3();  
do_stuff2();  
cilk_sync;  
do_stuff4();  
...
```


Another Execution DAG


- DAG represents the series-parallel structure of the execution of a Cilk Plus program
- Example:
 - two spawns (A) & (B)
 - one sync (C)




Note: computation on edges

Work and Span

- Edges represent serial computation (work)



- Span: most expensive path from beginning to end
—also known as critical path length



Note: computation on edges

cilkview

- **Rewrites executable to measure execution in terms of work and span**
 - measures
 - **work** - total # instructions executed, w/o parallel ovhd
 - **span** - # instructions executed on the critical path (w/o ovhd)
 - **burdened span** - # instructions executed on critical path (incl ovhd)
 - **parallelism** - work/span (max speedup on infinite cores, w/o ovhd)
 - **burdened parallelism** - $\text{work}/(\text{burdened span})$
 - **number of spawns/syncs**
 - **average instructions per strand** - $\text{work}/\text{strands}$
 - **strands along span** - # strands in the critical path
 - **average instructions / strand on span** = $\text{work}/(\text{strands along span})$
 - **total number of atomic instructions** - e.g., used for locks
 - **frame count**
- **Predicts speedup on various numbers of processors based on work and span**

cilkview Demo

- Explore cilkview for performance analysis

- cp /projects/comp422/cilkplus/examples/fib ~/fib

- cd ~/fib

- cilkview ./fib 20

- cilkview ./fib 30

- cilkview ./fib 35

- cilkview ./fib-trunc 35 10

Remaining Cilk Plus Features

- Introduce Cilk Plus array notation
- Gain more experience with Cilk Plus reducers
 - `cp /projects/comp422/cilkplus/examples/features-tutorial .`
 - array_notations: vector notation in Cilk Plus
 - reducers: more reducer examples
 - each directory contains a Makefile that can build and run all examples

Cilk Plus Array Notation

- **Elementwise arithmetic**

`c[:] = a[:] + 5;`

- **Set even rows in a 2D array**

`b[0:5:2][:] = 12;`

- **Vector conditionals**

`// Check and report each element containing 5 w/ Array Notation`

`if (5 == a[:]) an_results[:] = "Matched";`

`else an_results[:] = "Not Matched";`

- **Vector conditionals**

`// Call a fn on each element of a vector using Array Notation`

`fn(a[:]);`

– On STIC,
see </projects/comp422/cilkplus/examples/features-tutorial>

Recall: Task Scheduling in Cilk

Strategies

- **Work-stealing:** processor looks for work when it becomes idle
- **Lazy parallelism:** don't realize parallelism until necessary
 - benefits:
 - executes with precisely as much parallelism as needed
 - minimizes the number of threads that must be set up
 - runs with same efficiency as serial program on uniprocessor

Compilation Strategy

MIT Cilk generates two copies of each procedure

- **Fast clone: for optimized execution on a single processor**
 - spawned threads are fast
- **Slow clone: triggered by work stealing, full parallel support**
 - used to handle execution of “stolen procedure frames”
 - supports Cilk’s work-stealing scheduler
 - few steals when enough parallel slackness exists
 - speed of slow copy is not critical for performance
- **“Work-first” principle: minimize cost in fast clone**

Two Schedulers

- **Nanoscheduler: compiled into cilk program**
 - execute cilk function and spawns in exactly the same order as C
 - on one PE: when no microscheduling needed, same order as C
 - efficient coordination with microscheduler
- **Microscheduler**
 - schedule procedures across a fixed set of processors
 - implementation: randomized work-stealing scheduler
 - when a processor runs out of work, it becomes a thief
 - steals from victim processor chosen uniformly at random

Nanscheduler Sketch

Fast clone

- Upon entering a **cilk** function
 - allocate a frame in the heap
 - initialize frame to hold function's state
 - push the frame on the bottom of a deque
 - frame on stack \leftrightarrow frame in deque
- At a **spawn**
 - save function state into the frame
 - only live, dirty variables
 - save the entry number into the frame
 - call spawned procedure as a function
- After each spawn
 - check to see if parent has been stolen
 - if frame is still in the deque, it has not
 - if so, clean up C stack
- Each **sync** becomes a no-op
- When the procedure returns
 - pop the frame off the deque
 - resume caller after the spawn

```
int fib (int n)
{
    fib_frame *f;
    f = alloc(sizeof(*f));
    f->sig = fib_sig;
    if (n<2) {
        free(f, sizeof(*f));
        return n;
    }
    else {
        int x, y;
        f->entry = 1;
        f->n = n;
        *T = f;
        push();
        x = fib (n-1);
        if (pop(x) == FAILURE)
            return 0;
        ...
        ;
        free(f, sizeof(*f));
        return (x+y);
    }
}
```

*frame pointer
allocate frame
initialize frame*

free frame

*save PC
save live vars
store frame pointer
push frame
do C call
pop frame
frame stolen
second spawn
sync is free!
free frame*

Fast Clone and Nanoscheduler

- **Fast clone is never stolen**
 - converted to slow when steal occurs
 - enables optimizations
- **No sync needed in fast clone**
 - no children have been spawned
- **Frame saves state:**
 - PC (entry number)
 - live, dirty variables
- **Push and pop must be fast**

Nanoscheduler Overheads

Basis for comparison: serial C


- Allocation and initialization of frame, push onto 'stack'
—a few assembly instructions
- Procedure's state needs to be saved before each spawn
—entry number, live variables
- Check whether frame is stolen after each spawn
—two reads, compare, branch
- On return, free frame - a few instructions
- One extra variable to hold frame pointer

Runtime Support for Scheduling

Each processor has a ready deque (doubly ended queue)

- **Tail:** worker adds or removes procedures (like C call stack)
- **Head:** thief steals from head of a victim's deque

Deque for a Process




- ***Deque* grows downward**
- ***Stack frame* contains local variables for a procedure invocation**
- **Procedure call → new frame is pushed onto the bottom of the deque**
- **Procedure return → bottom frame is popped from the deque**
- **Deque maintains order (synchronizes) between caller and callee**

Cilk's Cactus Stacks


A cactus stack enables sharing of a C function's local variables

```
void A() { B(); C(); }  
void B() { D(); E(); }  
void C() { F(); }  
void D() {}  
void E() {}  
void F() {}
```

call tree



each procedure's view of stack



Rules

- pointers can be passed down call chain
- only pass pointers up if they point to heap
 - functions cannot return ptrs to local variables

Microscheduler

Schedule procedures across a fixed set of processors

- When a processor runs out of work, it becomes a **thief**
 - steals from **victim** processor chosen uniformly at random
- When it finds victim with frames in its deque
 - takes the topmost frame (least recently pushed)
 - places frame into its own deque
 - gives the corresponding procedure to its own nanoscheduler
- Microscheduler executes slow clone
 - receives only pointer to frame as argument
 - real args and local state in frame
 - restores pgm counter to proper place using switch stmt (Duff's device)
 - at a **sync**, must wait for children
 - before the procedure returns, place return value into frame

Coordinating Thief and Worker

Options

- Always use a lock to manipulate each worker's deque
- Use protocol that only relies on atomicity of read and write
—based on ideas from a locking protocol by Dijkstra

Simplified THE Protocol (Without the 'E')

- **Shared memory deque**
 - **T: first unused**
 - **H: head**
 - **E: exception**
- **Work-first**
 - **move costs from worker to thief**
- **One worker per deque**
- **One thief at a time**
 - **enforced by lock**

Worker/Victim

```
1  push() {
2      T++;
3  }
4  pop() {
5      T--;
6      if (H > T) {
7          T++;
8          lock(L);
9          T--;
10         if (H > T) {
11             T++;
12             unlock(L);
13             return FAILURE;
14         }
15         unlock(L);
16     }
17     return SUCCESS;
18 }
```

Thief

```
1  steal() {
2      lock(L);
3      H++;
4      if (H > T) {
5          H--;
6          unlock(L);
7          return FAILURE;
8      }
9      unlock(L);
10     return SUCCESS;
11 }
```

- actions on tail contribute to work overhead
- actions on head contribute only to critical path overhead


Deque Pop

Three cases

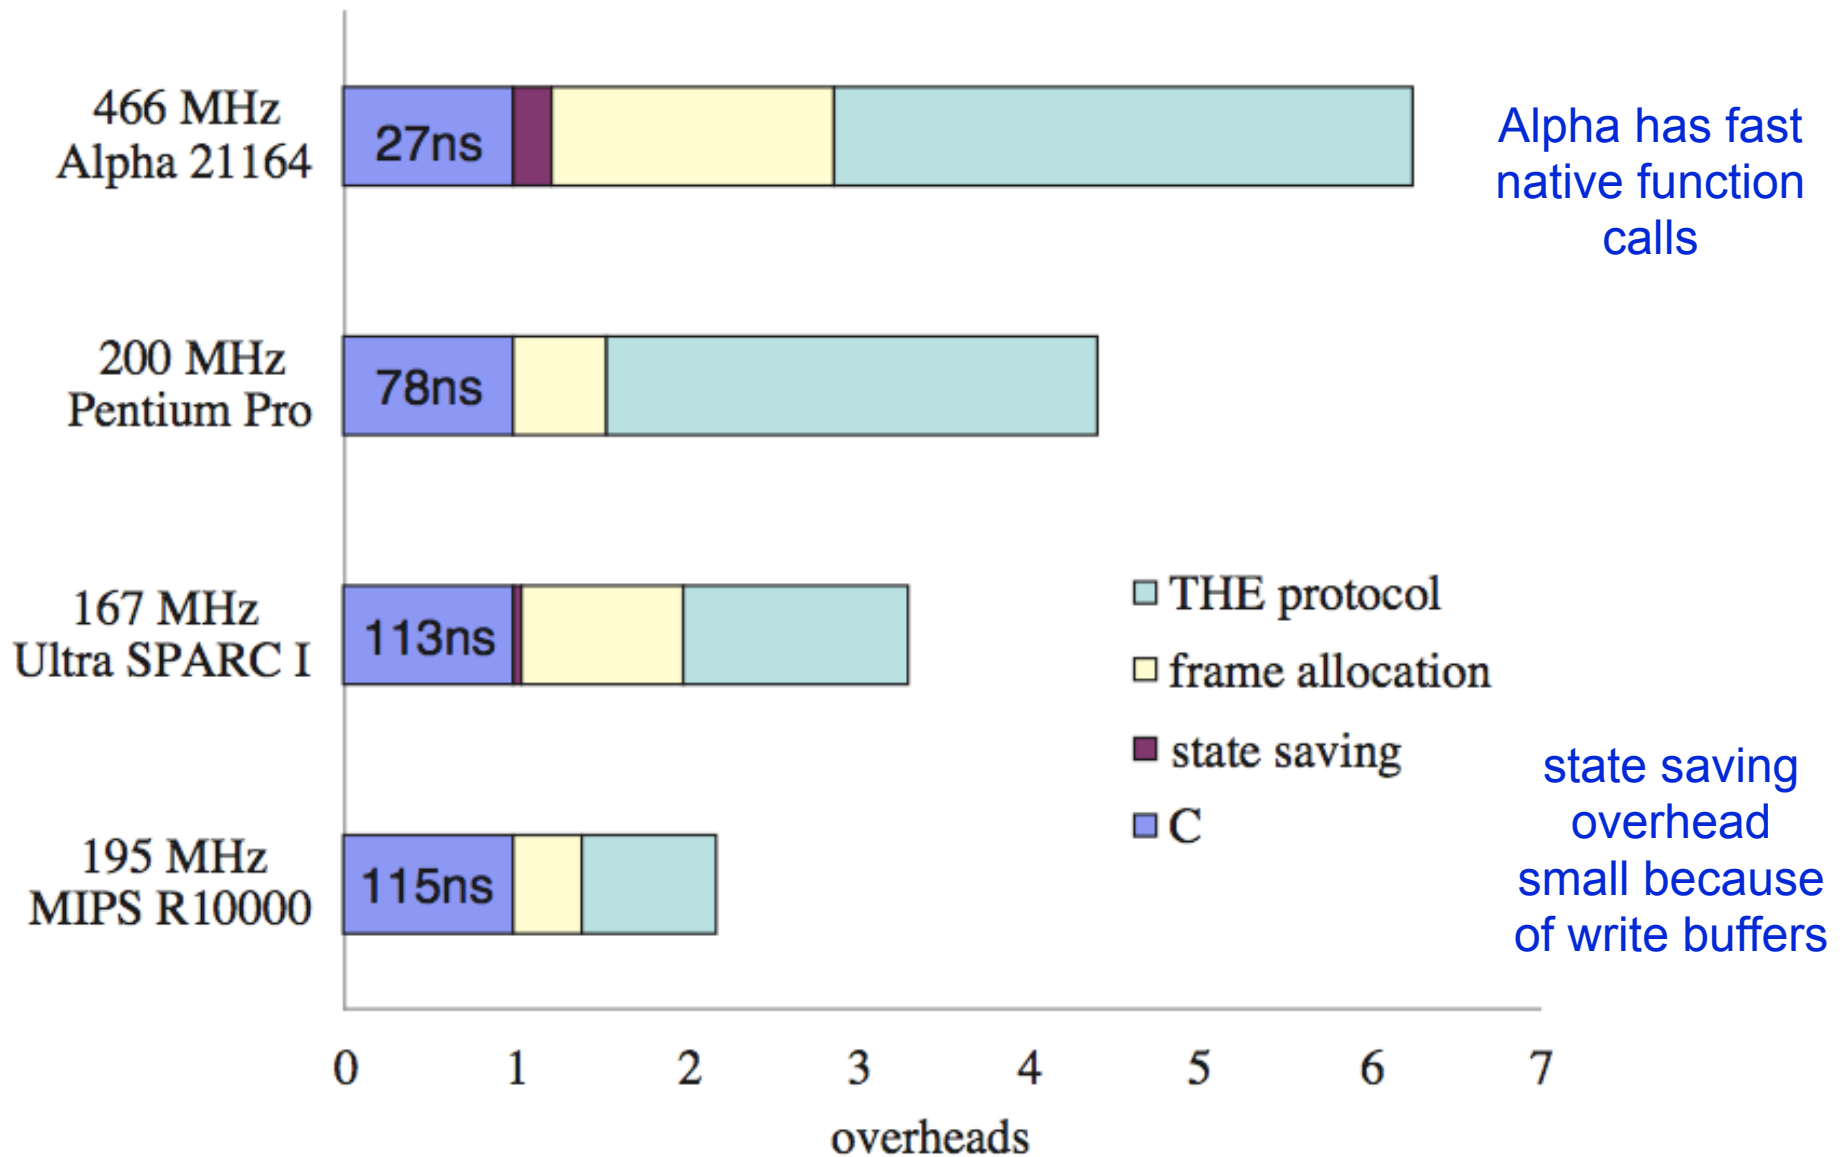
(a) no conflict

(b) At least one
(thief or victim)
finds ($H > T$) and
backs up; other
succeeds

(c) Deque is empty,
both threads
return



Work Overhead for **fib**



References

- **Charles E. Leiserson. Cilk LECTURE 1. Supercomputing Technologies Research Group. Computer Science and Artificial Intelligence Laboratory. <http://bit.ly/mit-cilk-lec1>**
- **Charles Leiserson, Bradley Kuzmaul, Michael Bender, and Hua-wen Jing. MIT 6.895 lecture notes - Theory of Parallel Systems. <http://bit.ly/mit-6895-fall03>**
- **Matteo Frigo, Charles Leiserson, and Keith Randall. The implementation of the Cilk-5 multithreaded language. In PLDI (Montreal, Quebec, Canada, June 17 - 19, 1998), 212-223.**
- **Intel Cilk++ Programmer's Guide. Document # 322581-001US.**
- **Matteo Frigo, Pablo Halpern, Charles E. Leiserson, and Stephen Lewin-Berlin. Reducers and other Cilk++ hyperobjects. SPAA '09, 79-90. Talk Slides. April 11, 2009. <http://bit.ly/reducers>**