

---

# Unified Parallel C (UPC)


**John Mellor-Crummey**

**Department of Computer Science  
Rice University**


**`johnmc@rice.edu`**

# Idealized Parallel Architectures

---



Shared Memory




Distributed Memory

## Programming Models


Cilk  
OpenMP  
Pthreads

MPI  
UPC  
CAF



# Idealized Parallel Architectures of Today

---



Hybrid Shared + Distributed Memory

Programming Models

e.g., MPI + OpenMP  
PGAS models


# Performance Concerns for Distributed Memory

---

**Data movement and synchronization are expensive**

**To minimize overheads**

- **Co-locate data with processes**
- **Aggregate multiple accesses to remote data**
- **Overlap communication with computation**



**Distributed Memory**

# Partitioned Global Address Space Languages

---

- **Global address space**
  - one-sided communication (GET/PUT) simpler than msg passing
- **Programmer has control over performance-critical factors**
  - data distribution and locality control lacking in OpenMP
  - computation partitioning mostly up to the compiler
  - communication placement } & runtime system for OpenMP and Cilk
- **Data movement and synchronization as language primitives**
  - amenable to compiler-based communication optimization

# Partitioned Global Address Space Languages

---

- Unified Parallel C (C) <http://upc.wikinet.org>
- Titanium (Java) <http://titanium.cs.berkeley.edu>
- Coarray Fortran 2.0 (Fortran) <http://caf.rice.edu>
- Related efforts: HPCS Languages
  - X10 (<http://x10-lang.org>)
  - Chapel (<http://chapel.cray.com>)
  - Fortress (<https://projectfortress.java.net>)

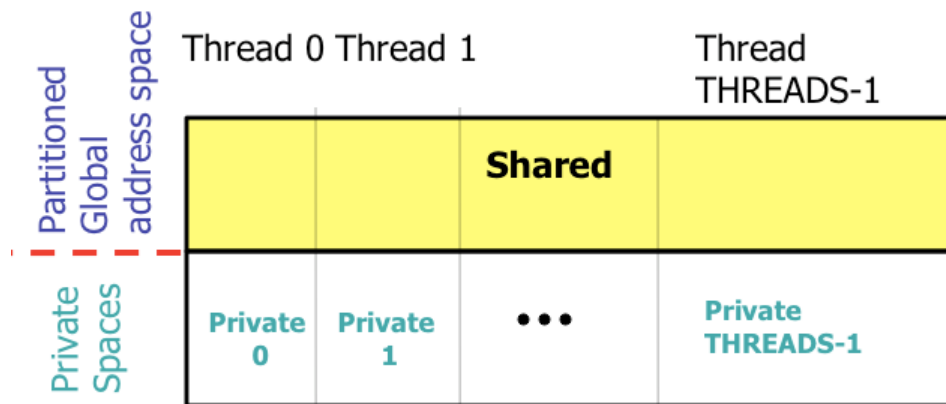
# Unified Parallel C (UPC)

---

- An explicit parallel extension of the C language
  - a few extra keywords
    - `shared`, `MYTHREAD`, `THREADS`, `upc_forall`
- Language features
  - partitioned global address space for shared data
    - part of shared data co-located with each thread
  - threads created at application launch
    - each bound to a hardware thread / core
    - each has some private data
  - a memory model
    - defines semantics of interleaved accesses to shared data
  - synchronization primitives
    - barriers
    - locks
    - load/store

# UPC Execution Model

- Multiple threads work independently in a SPMD fashion
  - MYTHREAD specifies thread index (0..THREADS-1)
  - # threads specified at compile-time or program launch
- Address Space



- Threads synchronize as necessary using
  - synchronization primitives
  - shared variables



# Shared and Private Data

---

- Static and dynamic memory allocation of each type of data
- Shared objects placed in memory based on affinity
  - shared scalars have affinity to thread 0
    - here, a scalar means a singleton instance of any type
  - elements of shared arrays are allocated round robin among memory modules co-located with each thread

# A One-dimensional Shared Array

---

Consider the following data layout directive

```
shared int y[2 * THREADS + 1];
```

For  $\text{THREADS} = 3$ , we get the following layout

Thread 0

y[0]
y[3]
y[6]

Thread 1

y[1]
y[4]

Thread 2

y[2]
y[5]

# A Multi-dimensional Shared Array

---

```
shared int A[4][THREADS];
```

For THREADS = 3, we get the following layout

Thread 0

A[0][0]
A[1][0]
A[2][0]
A[3][0]

Thread 1

A[0][1]
A[1][1]
A[2][1]
A[3][1]

Thread 2

A[0][2]
A[1][2]
A[2][2]
A[3][2]


# Shared and Private Data

---

Consider the following data layout directives


```
shared int x; // x has affinity to thread 0
shared int y[THREADS];
int z;        // private
```

For THREADS = 3, we get the following layout



# Controlling the Layout of Shared Arrays

- Can specify a blocking factor for shared arrays
  - default block size is 1 element
- Shared arrays are distributed on a block per thread basis, round robin allocation of block size chunks
- Example layout using block size specifications
  - e.g., `shared [2] int a[16]` block size



# Blocking of Shared Arrays

---


- Block size and THREADS determine *affinity*  
—with which thread will a datum be co-located
- Element  $i$  of a blocked array has affinity to thread:

$$\left\lfloor \frac{i}{\text{blocksize}} \right\rfloor \bmod \text{THREADS}$$


# Blocking Multi-dimensional Data I

- Manage the interaction between
  - contiguous memory layout of C multi-dimensional arrays
  - blocking factor for shared layout
- Consider layouts for different block sizes for
  - shared **[BLOCKSIZE]** double grids[N][N];


For the case where  $N = K * \text{THREADS}$ :



Default  
 $\text{BLOCKSIZE}=1$



Column Blocks  
 $\text{BLOCKSIZE}=N/\text{THREADS}$



Distribution by Row  
 $\text{BLOCKSIZE}=N$

# Blocking Multi-dimensional Data II

- Consider the data declaration

—shared [3] int A[4][THREADS];

- When THREADS = 4, this results in the following data layout

Thread 0	Thread 1	Thread 2	Thread 3
A[0][0]	A[0][3]	A[1][2]	A[2][1]
A[0][1]	A[1][0]	A[1][3]	A[2][2]
A[0][2]	A[1][1]	A[2][0]	A[2][3]
A[3][0]	A[3][3]		
A[3][1]			
A[3][2]			

The mapping is not pretty when the rightmost dimensions aren't a multiple of THREADS



# A Simple UPC Program: Vector Addition

```
//vect_add.c
#include <upc_relaxed.h>
#define N 100*THREADS

shared int v1[N], v2[N], v1plusv2[N];

void main() {
    int i;
    for(i=0; i<N; i++)
        if (MYTHREAD == i % THREADS)
            v1plusv2[i]=v1[i]+v2[i];
}
```

Iteration #:

Thread 0 Thread 1

0  
2

1  
3

v1[0]  
v1[2]

v1[1]  
v1[3]

...

v2[0]  
v2[2]

v2[1]  
v2[3]

...

v1plusv2[0]

v1plusv2[1]

v1plusv2[2]

v1plusv2[3]

...

Shared Space

Each thread executes each iteration to check if it has work

# A More Efficient Vector Addition

```
//vect_add.c
#include <upc_relaxed.h>
#define N 100*THREADS

shared int v1[N], v2[N], v1plusv2[N];

void main() {
    int i;
    for(i = MYTHREAD; i < N; i += THREADS)
        v1plusv2[i]=v1[i]+v2[i];
}
```

Iteration #:

Thread 0 Thread 1

0  
2

1  
3

v1[0]	v1[1]
v1[2]	v1[3]

...

v2[0]	v2[1]
v2[2]	v2[3]

...

v1plusv2[0]	v1plusv2[1]
v1plusv2[2]	v1plusv2[3]

...

Shared Space

Each thread executes only its own iterations

# Worksharing with `upc_forall`

---

- Distributes independent iterations across threads
- Simple C-like syntax and semantics
  - `upc_forall`(init; test; loop; affinity)
- Affinity is used to enable locality control
  - usually, map iteration to thread where the iteration's data resides
- Affinity can be
  - an integer expression, or a
  - reference to (address of) a shared object

# Work Sharing + Affinity with `upc_forall`

---

- Example 1: explicit affinity using shared references

```
shared int a[100], b[100], c[100];  
int i;  
upc_forall (i=0; i<100; i++; &a[i])  
    a[i] = b[i] * c[i];
```

- Example 2: implicit affinity with integer expressions

```
shared int a[100], b[100], c[100];  
int i;  
upc_forall (i=0; i<100; i++; i)  
    a[i] = b[i] * c[i];
```

**Note:** both yield a round-robin distribution of iterations

# Vector Addition Using `upc_forall`

thread affinity for work: have  
thread  $i$  execute iteration  $i$

```
//vect_add.c
#include <upc_relaxed.h>
#define N 100*THREADS

shared int v1[N], v2[N], v1plusv2[N];

void main()
{
    int i;
    upc_forall(i = 0; i < N; i++; i)
        v1plusv2[i] = v1[i] + v2[i];
}
```

Iteration #:

Thread 0	Thread 1
0	1
2	3
v1[0]	v1[1]
v1[2]	v1[3]
...	
v2[0]	v2[1]
v2[2]	v2[3]
...	
v1plusv2[0]	v1plusv2[1]
v1plusv2[2]	v1plusv2[3]
...	

Shared Space

Each thread executes subset of global  
iteration space as directed by affinity clause

# Work Sharing + Affinity with `upc_forall`

---

- **Example 3: implicit affinity by chunks**

```
shared int a[100], b[100], c[100];  
int i;  
upc_forall (i=0; i<100; i++; (i*THREADS)/100)  
    a[i] = b[i] * c[i];
```


- **Assuming 4 threads, the following results**

i	i*THREADS	i*THREADS/100
0..24	0..96	0
25..49	100..196	1
50..74	200..296	2
75..99	300..396	3

# Matrix-Vector Multiply (Default Distribution)

```
// vect_mat_mult.c
#include <upc_relaxed.h>


shared int a[THREADS][THREADS];
shared int b[THREADS], c[THREADS];
void main (void) {
    int i, j;
    upc_forall (i = 0; i < THREADS; i++; i) {
        c[i] = 0;
        for ( j= 0 ; j < THREADS; j++)
            c[i] += a[i][j]*b[j];
    }
}
```



# Matrix-Vector Multiply (Better Distribution)

```
// vect_mat_mult.c
#include <upc_relaxed.h>

shared [THREADS] int a[THREADS][THREADS];
shared int b[THREADS], c[THREADS];
void main (void) {
    int i, j;
    upc_forall( i = 0 ; i < THREADS ; i++; i) {
        c[i] = 0;
        for ( j= 0 ; j< THREADS ; j++)
            c[i] += a[i][j]*b[j];
    }
}
```





# UPC Pointers

- Needed for expressive data structures

- Flavors

—private pointers  
pointing to local

– `int *p1`

—private pointers  
pointing to shared

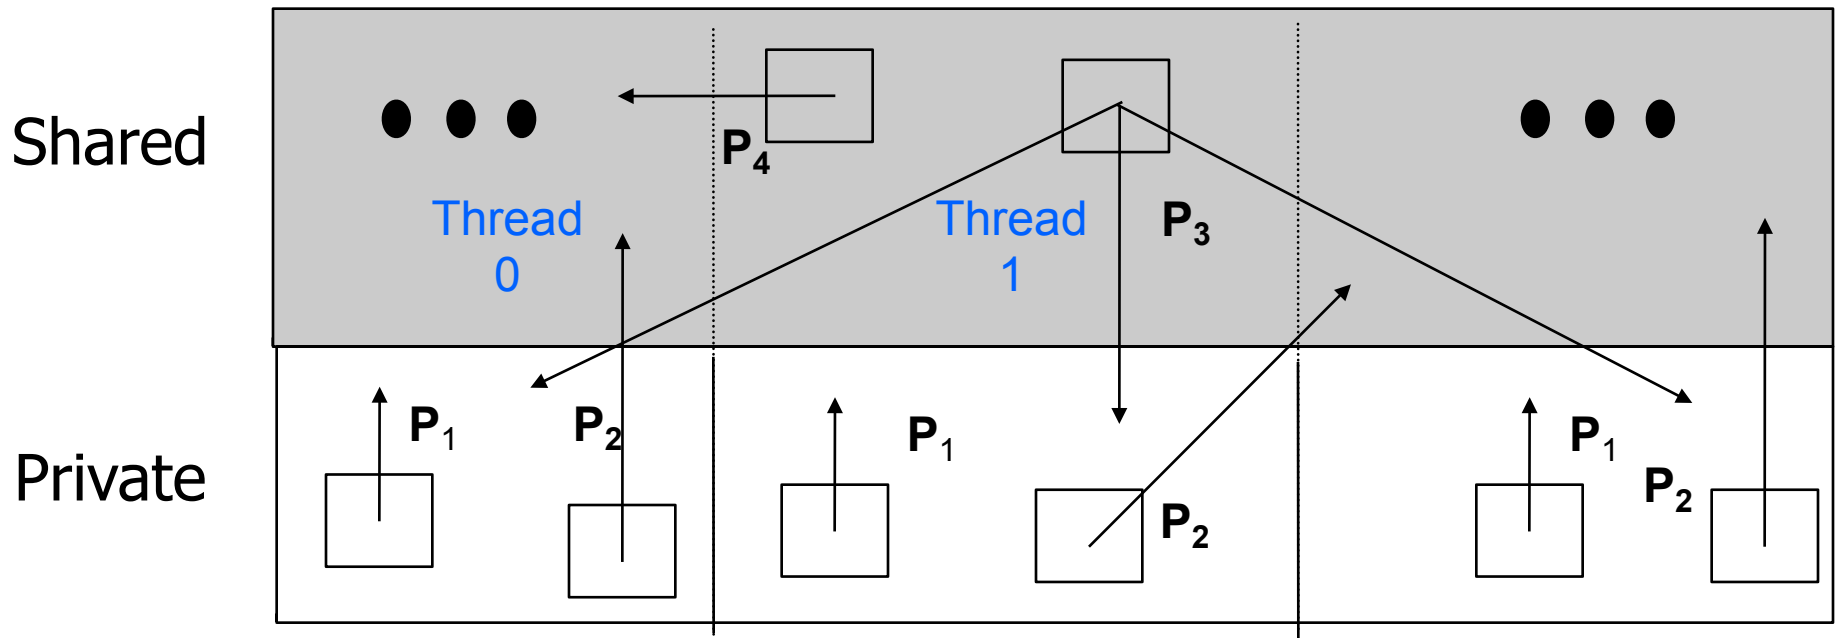
– `shared int *p2`

—shared pointers  
pointing to local

– `int *shared p3`

—shared pointers  
pointing to shared

– `shared int *shared p4`



# UPC Pointer Implementation Requirements

---

- **Handle shared data**
- **Support pointer arithmetic**
- **Support pointer casting**

# UPC Pointer Representation

- **UPC pointers to shared objects have three fields**
  - thread number
  - local address of block
  - phase (specifies position in the block)

Thread #	Block Address	Phase
----------	---------------	-------

- **Example: Cray T3E implementation**

Phase		Thread	Virtual Address	
63	49	48	38	37
				0

# UPC Pointer Features

---

- **Pointer arithmetic**
  - supports blocked and non-blocked array distributions
- **Casting of shared to private pointers is allowed**
  - but not vice versa!
- **When casting a pointer-to-shared to a private pointer, the thread # of pointer-to-shared may be lost**
- **Casting of a pointer-to-shared to a private pointer**
  - well defined only if the target object has affinity to local thread

# Dynamic Memory Allocation of Shared Memory

---

- **Dynamic memory allocation of shared memory is available**
- **Functions can be collective or not**
- **Collective function**
  - called by every thread
  - returns the same value to each of them
- **Collective function names typically include “all”**

# Global Allocation of Shared Memory

---

- **Collective allocation function: must be called by all threads**

- shared void \*upc\_all\_alloc(size\_t nblocks, size\_t nbytes)

- nblocks: number of blocks

- nbytes: block size

- each thread gets the same pointer

- equivalent to :

- shared [nbytes] char[nblocks \* nbytes]

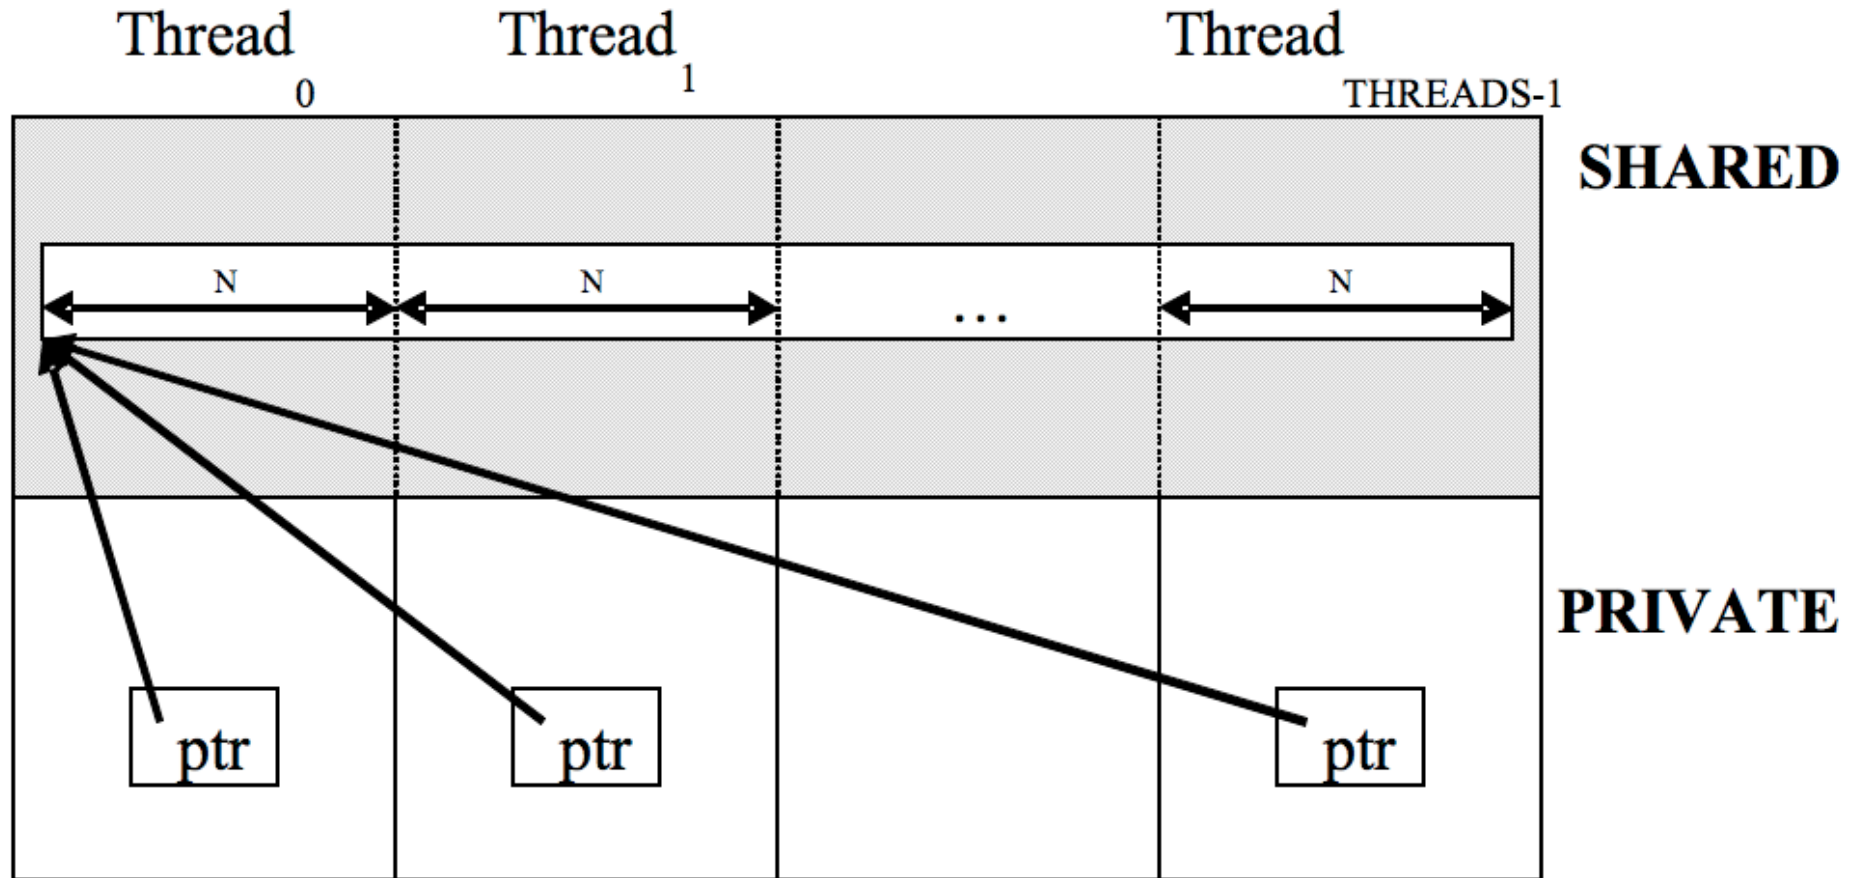
- **Non-collective version that yields the same layout**

- shared void \*upc\_global\_alloc(size\_t nblocks, size\_t nbytes)

- nblocks: number of blocks

- nbytes: block size

# Global Allocation of Shared Memory



```
shared [N] int *ptr;  
ptr = (shared [N] int *)  
      upc_all_alloc( THREADS, N*sizeof( int ) );
```

# Local-Shared Memory Allocation

---

```
shared void *upc_alloc (size_t nbytes)
```

—nbytes:            block size

- Non collective; called by one thread
- Calling thread allocates a contiguous memory region in its local shared space
- Space allocated per calling thread is equivalent to shared [] char[nbytes]
- If called by more than one thread, multiple regions are allocated and each calling thread gets a different pointer



# Synchronization

---

- No implicit synchronization among the threads
- UPC provides the following synchronization mechanisms:
  - barriers
  - locks

# Synchronization - Barriers

---

- **Barriers (blocking)**
  - `upc_barrier expr_opt;`
- **Split-phase barriers (non-blocking)**
  - `upc_notify expr_opt;`
  - `upc_wait expr_opt;`
    - **note: `upc_notify` is not blocking `upc_wait` is**

# Synchronization - Locks

---

- **Lock primitives**
  - void upc\_lock(upc\_lock\_t \*l)
  - int upc\_lock\_attempt(upc\_lock\_t \*l) // success returns 1
  - void upc\_unlock(upc\_lock\_t \*l)
- **Locks are allocated dynamically, and can be freed**
- **Locks are properly initialized after they are allocated**

# Dynamic Lock Allocation

---

- **Collective lock allocation (à la upc\_all\_alloc)**  
—`upc_lock_t * upc_all_lock_alloc(void);`
- **Global lock allocation (à la upc\_global\_alloc)**  
—`upc_lock_t * upc_global_lock_alloc(void);`
- **Lock deallocation**  
—`void upc_lock_free(upc_lock_t *ptr);`

# Memory Consistency Model

---

- Dictates the ordering of shared operations
  - when a change to a shared object by a thread becomes visible to others
- Consistency can be **strict** or **relaxed**
- Relaxed consistency model
  - compiler & runtime can reorder accesses to shared data
- Strict consistency model
  - enforce sequential ordering of operations on shared data
    - no operation on shared can begin before previous ones are done
    - changes become visible immediately

# Memory Consistency

---

- Default behavior can be altered for a variable definition in the declaration using:
  - Type qualifiers: strict & relaxed
- Default behavior can be altered for a statement or a block of statements using
  - #pragma upc strict
  - #pragma upc relaxed
- Precedence order for memory consistency specifications
  - 1.declarations
  - 2.pragmas
  - 3.program level

# Memory Consistency Example

---

```
strict shared int flag_ready = 0;
shared int result0, result1;

if (MYTHREAD==0){
    results0 = expression1;
    flag_ready=1; //if not strict, it could be
                  // switched with the above statement
} else if (MYTHREAD==1){
    while(!flag_ready); //Same note
    result1=expression2+results0;
}
```

- Could have used a barrier between the first and second statement in the if and the else code blocks
  - **expensive: affects all operations at all threads**
- Above works as an example of point to point synchronization

# Forcing Memory Consistency via `upc_fence`

---

- What is a memory fence?
  - all memory operations initiated before a fence operation must complete before the fence completes
- UPC provides a fence construct
  - syntax
    - `upc_fence;`
  - semantics
    - equivalent to a null strict reference



# Library Operations for Bulk Data

---

- No flexible way to initiate bulk transfer operations in UPC
- Rely on library operations for bulk data transfer and set
  - void **upc\_memcpy**(shared void \* restrict dst, shared const void \* restrict src, size\_t n)
  - void **upc\_memget**(void \* restrict dst, shared const void \* restrict src, size\_t n);
  - void **upc\_mempu**t(shared void \* restrict dst, const void \* restrict src, size\_t n);
  - void **upc\_memset**(shared void \*dst, int c, size\_t n);

# Explicit Non-blocking Data Movement

---

- **Get, Put, Copy**

- `upc_handle_t bupc_memcpy_async(shared void *dst, shared const void *src, size_t nbytes)`
- `upc_handle_t bupc_memget_async(void *dst, shared const void *src, size_t nbytes)`
- `upc_handle_t bupc_mempush_async(shared void *dst, const void *src, size_t nbytes)`
  - same args and semantics as blocking variants
  - `upc_handle_t`: opaque handle representing the operation initiated

- **Synchronize using one of two new functions**

- `void bupc_waitsync(upc_handle_t handle)`
  - blocking test for completion
- `int bupc_trysync(upc_handle_t handle)`
  - non-blocking test for completion

# Meraculous De Novo Genome Assembler in UPC

- Chop reads into k-mers that overlap by k-1
- Store k-mers in distributed hash table.  
(key,value) =  
(k-mer, 2-char fwd/bwd extension [AGCT]  
[AGCT])
- From selected k-mers, perform forward and reverse traversals to construct contigs




Fig. 1: Performance and strong scaling of our de Bruijn graph construction & traversal and k-mer analysis steps on Cray XC30 for the human genome. The top three timing curves are with respect to the first y-axis (left) whereas the parallel efficiency curve is with respect to the second y-axis (right). The x-axis uses a log scale.




Fig. 2: Meraculous assembly flow chart.




Fig. 5: Parallel de Bruijn graph construction.

# Ongoing Work in UPC

---

- **Interoperability support for multicore**
  - UPC always has run on clusters of shared memory
  - Desire to run UPC on 1-n nodes for small n, within MPI
- **Adding hierarchy to memory (socket, node, cab?,...)**
- **Use unique network features (atomics, DMA, etc.)**
- **PGAS/UPC for GPUs (P = CPU/GPU partition + nodes)**
  - First step: GASnet on CPU clusters
  - Can put/get from GPU or CPU to/from GPU or CPU
  - (Illusion is a single address space; ugly on current hw)
- **SPMD PGAS + dynamic load balancing**
  - Dynamic load balancing is mostly useful for load imbalance in the apps, and cost across nodes is high
  - So apps choose to add tasking on top of UPC
- **Teams and autotuned collectives in UPC**

# References

---

## Slides adapted from

Tarek El-Ghazawi, Steven Seidel. High Performance Parallel Programming with Unified Parallel C. SC05 Tutorial. <http://upc.gwu.edu/tutorials/UPC-SC05.pdf>

## Meraculous De Novo Genome Assembler

E. Georganas et al., Parallel De Bruijn Graph Construction and Traversal for De Novo Genome Assembly, SC14: International Conference for High Performance Computing, Networking, Storage and Analysis, pp.437-448, Nov. 2014 doi: 10.1109/SC.2014.41  
URL: [http://www.eecs.berkeley.edu/~egeor/sc14\\_genome.pdf](http://www.eecs.berkeley.edu/~egeor/sc14_genome.pdf)