
Parallel Computing Platforms: Control Structures and Memory Hierarchy

John Mellor-Crummey

**Department of Computer Science
Rice University**

`johnmc@rice.edu`

Topics for Today

- **SIMD and MIMD control structure**
- **Memory hierarchy and performance**

Parallel Computing Platforms

A parallel computing platform must specify


- concurrency = control structure**
- interaction between concurrent tasks = communication model**

Control Structure of Parallel Platforms

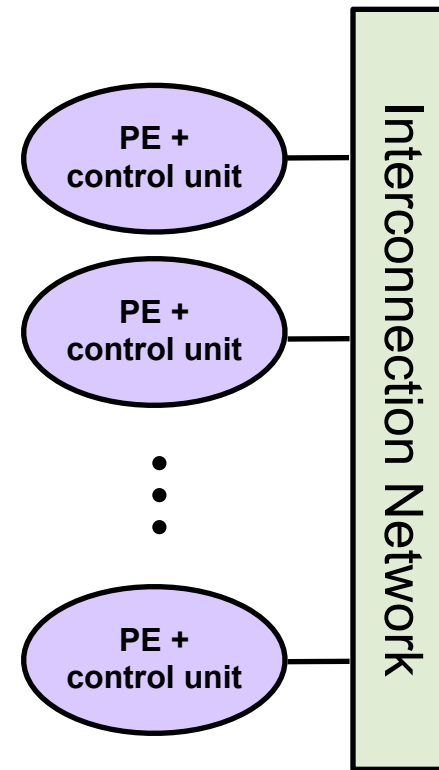
Parallelism ranges from instructions to processes

- **Processor control structure alternatives**
 - work independently
 - operate under the centralized control of a single control unit
- **MIMD**
 - Multiple Instruction streams**
 - each processor has its own control control unit
 - each processor can execute different instructions
 - Multiple Data streams**
 - processors work on their own data
- **SIMD (aka SIMT in nVIDIA GPUs)**
 - Single Instruction stream**
 - single control unit dispatches the same instruction to processors
 - Multiple Data streams**
 - processors work on their own data

SIMD and MIMD Processors



SIMD architecture



MIMD architecture


PE = Processing Element

SIMD Control

- **SIMD excels for computations with regular structure**
 - media processing, scientific kernels (e.g., linear algebra, FFT)
- **Activity mask**
 - per PE predicated execution: turn off operations on certain PEs
 - each PE tests own conditional and sets own activity mask
 - PE can conditionally perform operation predicated on mask value

Example: 128-bit SIMD Vectors

- Data types: anything that fits into 16 bytes, e.g.,



- Instructions operate in parallel on data in this 16 byte register
 - add, multiply etc.
- Data bytes must be **contiguous in memory** and **aligned**
- Additional instructions needed for
 - masking data
 - moving data from one part of a register to another

Computing with SIMD Vector Units

- **Scalar processing**
—one operation produces one result

X

+

Y

X + Y

- **SIMD vector units**
—one operation produces multiple results

X

x3 x2 x1 x0

+

Y

y3 y2 y1 y0

X + Y

x3+y3 x2+y2 x1+y1 x0+y0

Executing a Conditional on a SIMD Processor

conditional statement

```
if (A == 0)
  then C = B
  else C = B/A
```

initial values

A	0
B	5
C	0

Processor 0

A	4
B	8
C	0

Processor 1

A	2
B	2
C	0

Processor 2

A	0
B	7
C	0

Processor 3

execute
"then" branch

A	0
B	5
C	5

Processor 0

A	4
B	8
C	0

Processor 1

A	1
B	2
C	0

Processor 2

A	0
B	0
C	7

Processor 3

execute
"else" branch

A	0
B	5
C	5

Processor 0

A	4
B	8
C	2

Processor 1

A	2
B	2
C	1

Processor 2

A	0
B	0
C	7


Processor 3

SIMD Examples

- **Previously: SIMD computers**
 - e.g., **Connection Machine CM-1/2, and MasPar MP-1/2**
 - **CM-1 (1980s): 65,536 1-bit processors**
- **Today: SIMD units or co-processors (AKA accelerators)**
 - vector units**
 - **SSE/2/3/4 - Streaming SIMD Extensions**
 - 8 128-bit vector registers**
 - 128 bits as 8-bit chars, 16-bit words, 32/64-bit int and float
 - **AVX - Advanced Vector Extensions**
 - 16 256-bit vector registers in Intel and AMD processors since 2011**
 - 256 bits as 8-bit chars, 16-bit words, 32/64-bit int and float
 - 32 512-bit vector registers in Intel Xeon Phi**
 - 512 bits as 8-bit chars, 16-bit words, 32/64-bit int and float
 - co-processors**
 - **nVIDIA Kepler GK110 GPGPU**
 - **ClearSpeed CSX700 array processor (control PE + array of 96 PEs)**
 - <http://www.clearspeed.com>

Intel MIC / Xeon Phi 5110P (Nov 2012)

- 60 cores
 - 32KB L1 I/D cache
 - 512 KB L2
 - 32 512-bit vector registers
 - 4-way SMT per core
- 1.053 GHz
- 8GB GDDR5 memory
- Memory B/W 352 GB/s
- 1.011 TFLOP, 225 Watts

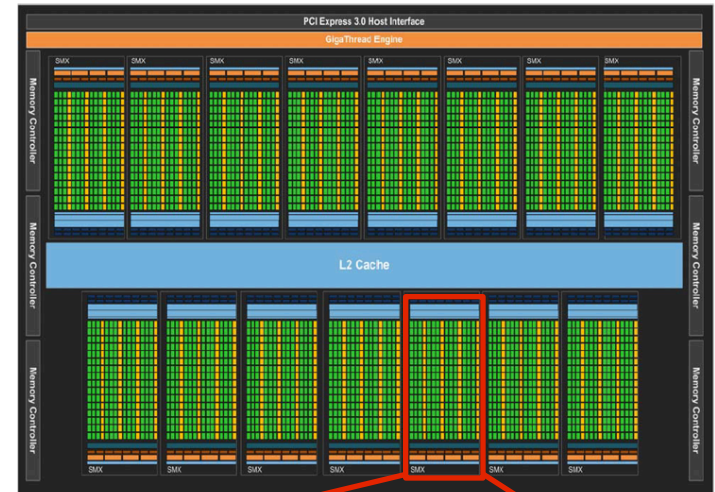


http://newsroom.intel.com/community/intel_newsroom/blog/2012/11/12/intel-delivers-new-architecture-for-discovery-with-intel-xeon-phi-coprocessors

NVIDIA Kepler GK110 (May 2012)

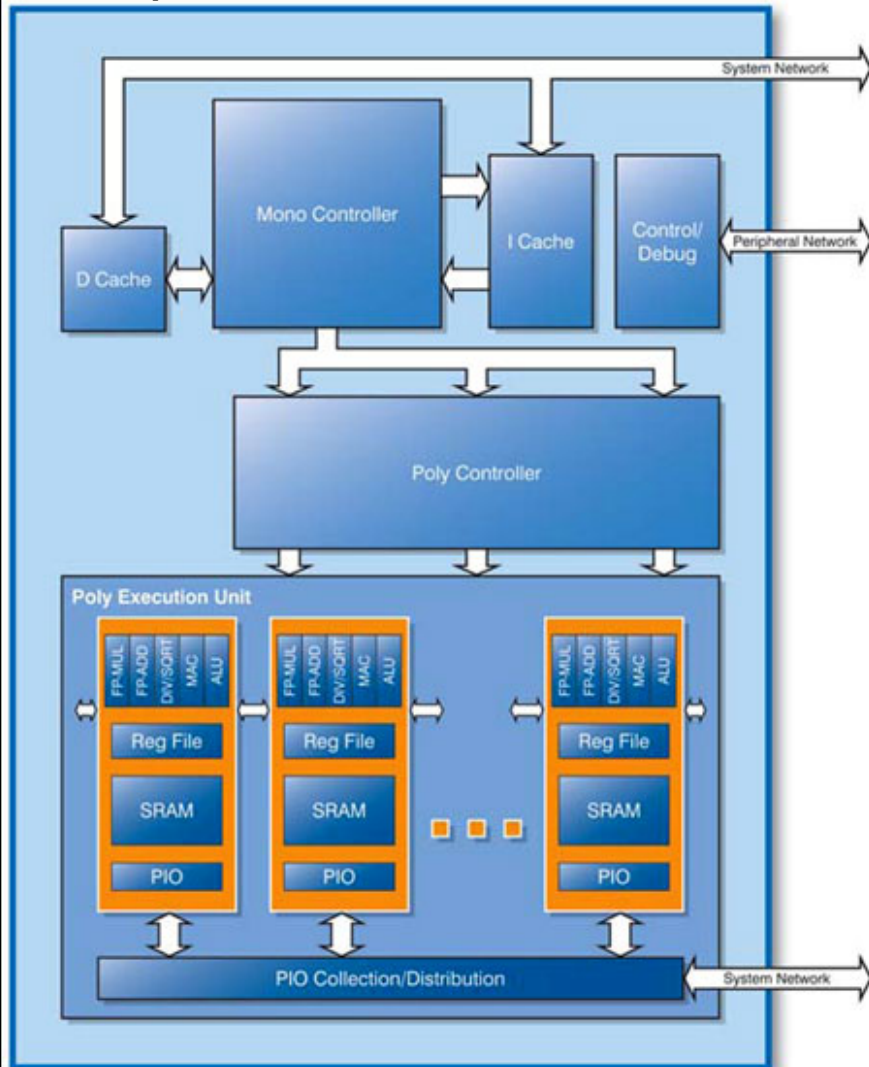
- 15 Streaming Multiprocessors (SMX)
- Each SMX
 - 192 CUDA cores (single precision)
 - fully pipelined FP and INT units
 - IEEE 754-2008; fused multiply add
 - four warp schedulers
 - 32-thread groups (warp)
 - 4 warps issue and execute concurrently
 - 2 inst/warp/cycle
 - 64 DP FP units
 - 32 LS units
 - 32 SFU
- 1.31 TFLOP, 235W

	KEPLER GK110
Compute Capability	3.5
Threads / Warp	32
Max Warps / Multiprocessor	64
Max Threads / Multiprocessor	2048
Max Thread Blocks / Multiprocessor	16
32-bit Registers / Multiprocessor	65536
Max Registers / Thread	255
Max Threads / Thread Block	1024
Shared Memory Size Configurations (bytes)	16K 32K 48K
Max X Grid Dimension	$2^{32}-1$
Hyper-Q	Yes
Dynamic Parallelism	Yes



SIMD: ClearSpeed MTAP Co-processor

MTAP processor



- **Features**
 - hardware multi-threading
 - asynchronous, overlapped I/O
 - extensible instruction set
- **SIMD core**
 - poly controller
 - poly execution unit
 - array of 192 PEs
 - 64- and 32-bit floating point
 - 250 MHz (key to low power)
 - 96 GFLOP, <15 Watts

Short Vectors: The Good and Bad

```
for (t = 0; t < T; ++t) {  
    for (i = 0; i < N; ++i)  
        for (j = 1; j < N+1; ++j)  
S1:      C[i][j] = A[i][j] + A[i][j-1];  
    for (i = 0; i < N; ++i)  
        for (j = 1; j < N+1; ++j)  
S2:      A[i][j] = C[i][j] + C[i][j-1];  
}
```

Performance:	AMD Phenom	1.2 GFlop/s
	Core2	3.5 GFlop/s
	Core i7	4.1 GFlop/s

(a) Stencil code


```
for (t = 0; t < T; ++t) {  
    for (i = 0; i < N; ++i)  
        for (j = 0; j < N; ++j)  
S3:      C[i][j] = A[i][j] + B[i][j];  
    for (i = 0; i < N; ++i)  
        for (j = 0; j < N; ++j)  
S4:      A[i][j] = B[i][j] + C[i][j];  
}
```

Performance:	AMD Phenom	1.9 GFlop/s
	Core2	6.0 GFlop/s
	Core i7	6.7 GFlop/s


(b) Non-Stencil code

The stencil code (a) has much lower performance than the non-stencil code (b) despite accessing 50% fewer data elements

- **Consider the following:**




Dimension-lifted Transformation (DLT)



**DLT +
vector
intrinsic**

**DLT +
autovec**

original



- (a) 1D array in memory
- (b) 2D view of same array
- (c) Transposed 2D array brings non-interacting elements into contiguous vectors
- (d) New 1D layout after transformation

Jacobi-1D:

$$a[i] = b[i-1] + b[i] + b[i+1]$$

MIMD Processors

Execute different programs on different processors

- **Platforms include current generation systems**
 - shared memory**
 - multicore laptop
 - workstation with multiple quad core processors
 - SGI Altix UV (up to 32K sockets, each with an 8-core processor)
 - Legacy: Cray X1 (up to 8K processors)
 - distributed memory**
 - clusters (e.g., biou.rice.edu, stic.rice.edu, davinci.rice.edu)
 - Cray XC, IBM Blue Gene
- **SPMD programming paradigm**
 - Single Program, Multiple Data streams**
 - same program on different PEs, behavior conditional on thread id**

SIMD vs. MIMD

- **SIMD platforms**

- special purpose: not well-suited for all applications
- custom designed with long design cycles
- less hardware: single control unit
- need less memory: only 1 copy of program
- today: SIMD common only for accelerators and vector units
 - accelerator: nVidia Kepler

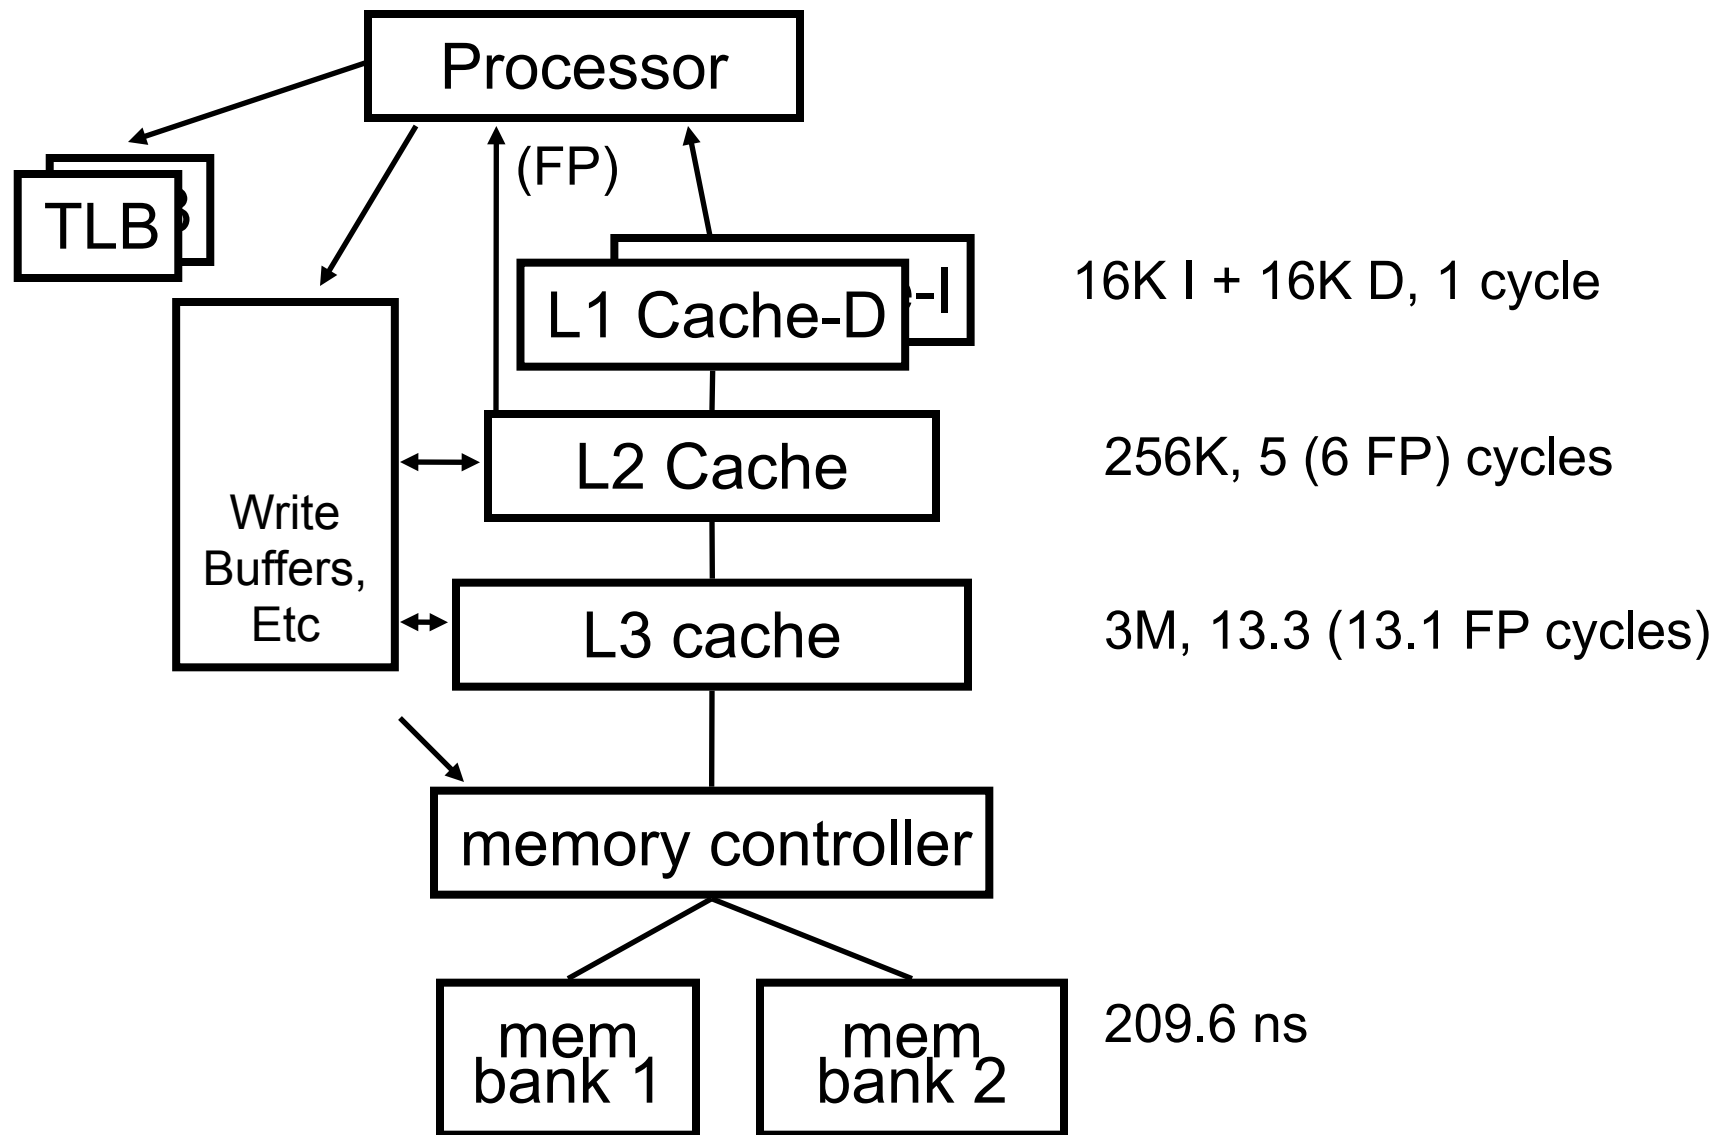
- **MIMD platforms**

- suitable for broad range of applications
- inexpensive: off-the-shelf components + short design cycle
- need more memory: program and OS on each processor

Data Movement and Communication

- **Latency**: How long does a single operation take?
 - measured in microseconds
- **Bandwidth**: What data rate can be sustained?
 - measured in Mbytes or GBytes per second
- These terms can be applied to
 - memory access
 - messaging

A Memory Hierarchy (Itanium 2)



Memory Bandwidth


- Limited by both
 - the bandwidth of the memory bus
 - the bandwidth of the memory modules
- Can be improved by increasing the size of memory blocks
- Memory system takes l time units to deliver b units of data
 - l is the latency of the system
 - b is the block size

Reusing Data in the Memory Hierarchy

- **Spatial reuse**: using more than one word in a multi-word line
—using multiple words in a cache line
- **Temporal reuse**: using a word repeatedly
—accessing the same word in a cache line more than once
- Applies at every level of the memory hierarchy
—e.g. TLB
 - spatial reuse: access multiple cache lines in a page
 - temporal reuse: access data on the same page repeatedly

Experimental Study of Memory (membench)


Microbenchmark for memory system performance



for array A of length L from 4KB to 8MB by 2x
for stride s from 4 Bytes (1 word) to L/2 by 2x
time the following loop
(repeat many times and average)
for i from 0 to L by **s**
load A[i] from memory (4 Bytes)

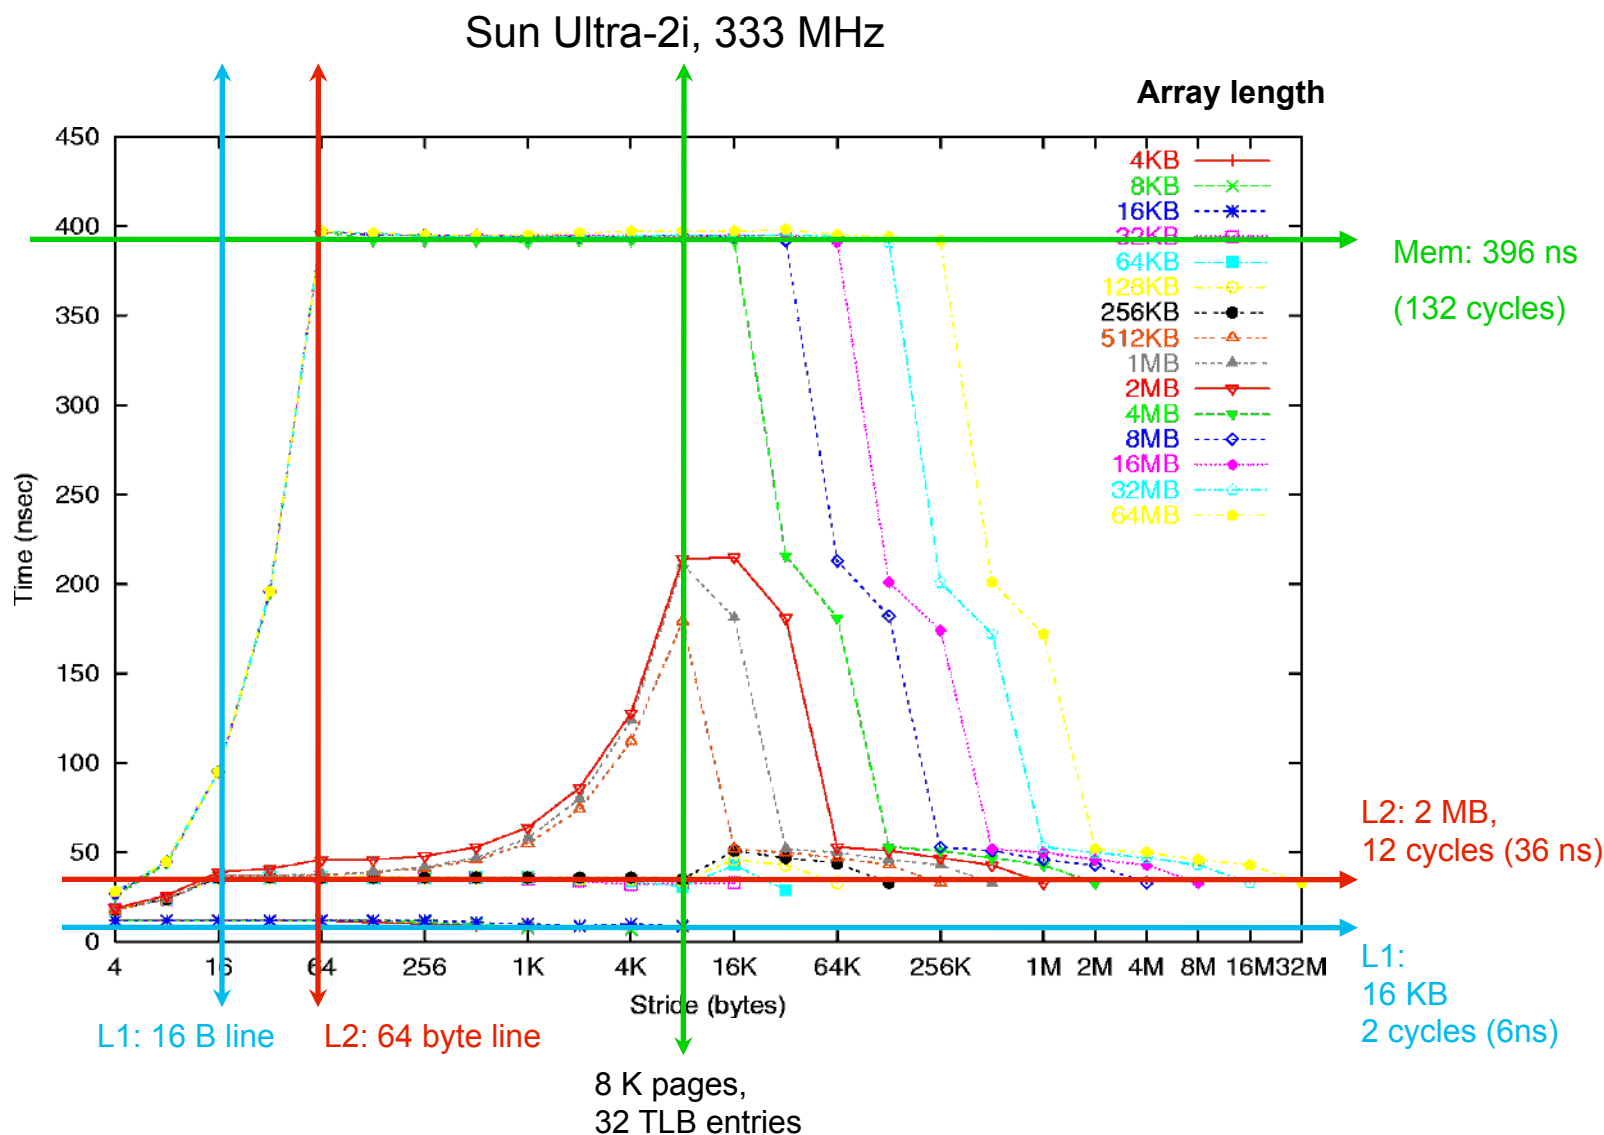
1 experiment

Membench: What to Expect




- Consider the average cost per load
 - plot one line for each array length, time vs. stride
 - unit stride is best: if cache line holds 4 words, only $\frac{1}{4}$ miss
 - if array is smaller than a cache, all accesses will hit after first run
 - time for first run is negligible with enough repetitions
 - upper right figure assumes only one level of cache
 - performance profile is more complicated on modern systems

Memory Hierarchy on a Sun Ultra-2i



See www.cs.berkeley.edu/~yelick/arvindk/t3d-isca95.ps for details

Memory Hierarchy on a Pentium III



Memory Bandwidth in Practice

What matters for application performance is “balance” between sustainable memory bandwidth and peak double-precision floating-point performance.

Systems at TACC

—Ranger (4-socket quad-core AMD “Barcelona”)

- bandwidth = 7.5 GB/s (2.19 GW/s, 8-Byte Words) per node
- peak FP rate = $2.3 \text{ GHz} * 4 \text{ FP Ops/Hz/core} * 4 \text{ cores/socket} * 4 \text{ sockets} = 147.2 \text{ GFLOPS/node}$
- ratio = 67 FLOPS/Word

—Lonestar (2-socket 6-core Intel “Westmere”)

- bandwidth = 41 GB/s (5.125 GW/s) per node
- peak FP rate = $3.33 \text{ GHz} * 4 \text{ Ops/Hz/core} * 6 \text{ cores/socket} * 2 \text{ sockets} = 160 \text{ GFLOPS/node}$
- ratio = 31 FLOPS/Word

—Stampede (2-socket 8-core Intel “Sandy Bridge” processors)

- bandwidth = 78 GB/s (9.75 GW/s) per node
- peak FP rate = $2.7 \text{ GHz} * 8 \text{ FP Ops/Hz} * 8 \text{ cores/socket} * 2 \text{ sockets} = 345.6 \text{ GFLOPS per node}$
- ratio = 35 FLOPS/Word

Credit: John McCalpin, TACC, <http://blogs.utexas.edu/jdm4372/2012/11/>

Memory System Performance: Summary

- Exploiting spatial and temporal locality is critical for
 - amortizing memory latency
 - increasing effective memory bandwidth
- Ratio # operations / # memory accesses
 - good indicator of anticipated tolerance to memory bandwidth
- Memory layout and computation organization significantly affect spatial and temporal locality

Multithreading for Latency Hiding

- A thread is a single stream of control in the flow of a program.
- We illustrate threads with a dense matrix vector multiply

```
for (i = 0; i < n; i++)  
    c[i] = dot_product(get_row(a, i), b);
```

- Each dot-product is independent of others
—thus, can execute concurrently
- Can rewrite the above code segment using threads

```
for (i = 0; i < n; i++)  
    c[i] = create_thread(dot_product, get_row(a, i), b);
```

Multithreading for Latency Hiding (contd)

- **Consider how the code executes**
 - first thread accesses a pair of vector elements and waits for them
 - second thread can access two other vector elements in the next cycle
 - ...
- **After l units of time**
 - (l is the latency of the memory system)
 - first thread gets its data from memory and performs its madd
- **Next cycle**
 - data items for the next function instance arrive
- ...
- **Every clock cycle, we can perform a computation**

Multithreading for Latency Hiding (contd)

- Previous example makes two hardware assumptions
 - memory system can service multiple outstanding requests
 - processor is capable of switching threads at every cycle
- Also requires program to have explicit threaded concurrency
- Machines such as the Sun T2000 (Niagara-2) and the Cray Threadstorm rely on multithreaded processors
 - can switch the context of execution in every cycle
 - are able to hide latency effectively
- Sun T2000, 64-bit SPARC v9 processor @1200MHz
 - organization: 8 cores, 4 strands per core, 8KB Data cache and 16KB Instruction cache per core, L2 cache: unified 12-way 3MB, RAM: 32GB
- Cray Threadstorm: 128 threads


Prefetching for Latency Hiding

- Misses on loads cause programs to stall; why not load data before it is needed?
 - by the time it is actually needed, it will be there!
- Drawback: need space to store early loads
 - may overwrite other necessary data in cache
 - if early loads are overwritten, we are little worse than before!
- Prefetching support
 - software only, e.g. Itanium2
 - hardware and software, e.g. Opteron
- Hardware prefetching requires
 - predictable access pattern
 - limited number of independent streams

Tradeoffs in Multithreading and Prefetching

- **Multithreaded systems**
 - bandwidth requirements
 - may increase very significantly because of reduced cache/ thread
 - can become bandwidth bound instead of latency bound
- **Multithreading and prefetching**
 - only address latency
 - may often exacerbate bandwidth needs
 - have significantly larger data footprint; need hardware for that

Understanding Performance Limitations



Williams, Waterman, Patterson; CACM April 2009


Shared Memory Platforms

- Part (or all) memory is accessible to all processors
 - modify data objects stored in shared memory
- Flavors of shared memory platforms
 - UMA: uniform memory access
 - time taken by a processor to access any memory word is identical
 - NUMA: non-uniform memory access
 - time taken by a processor to memory may vary


NUMA and UMA Platforms

- **NUMA and UMA platforms differ with respect to locality**
 - algorithms must exploit locality for performance on NUMA platforms
- **Programming these platforms**
 - easy communication: reads and writes are implicitly visible to other processors
 - tricky coordination: read-write operations to shared data must be coordinated
 - Pthreads mutexes, OpenMP critical sections
- **Cache coherent vs. non cache-coherent architectures**
 - non-cache coherent shared-address space architectures
 - provides an address map, but not coordinated sharing
 - processors must explicitly flush data to memory before sharing
 - examples: BBN Butterfly, Cray T3E
 - cache coherent architectures: caches coordinate access to multiple copies
 - hardware support to keep copies consistent
 - up to date values can be retrieved from cache by remote processors
 - examples: SGI Origin, SGI Altix


NUMA and UMA Platforms



UMA shared address
space platform with cache
(Sequent Symmetry, 1988)

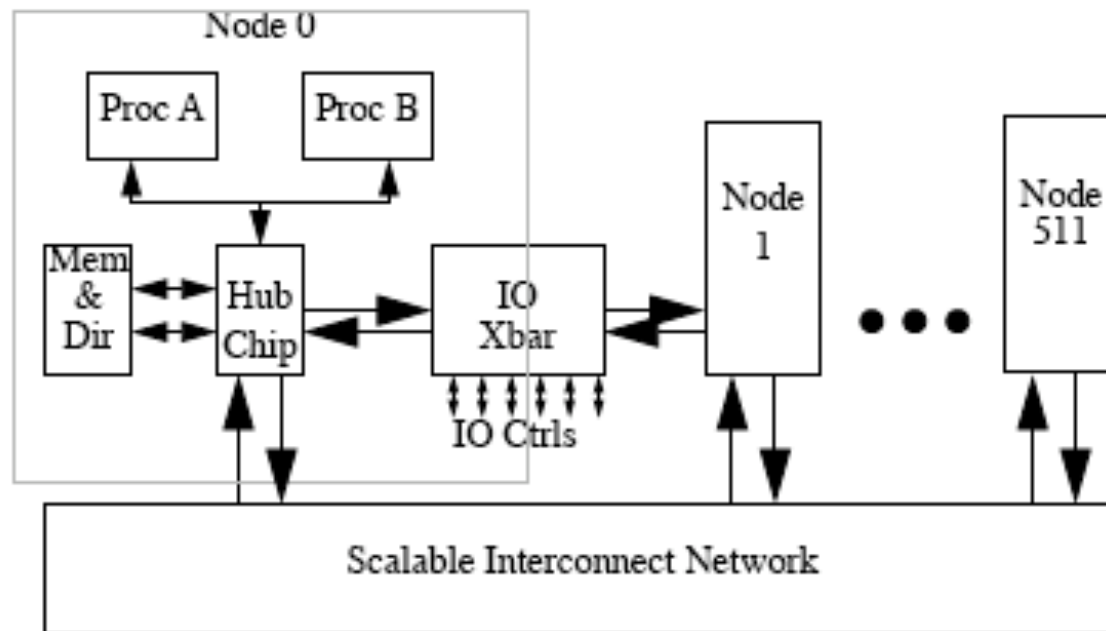


NUMA shared address
space platform
(BBN Butterfly, 1988)



UMA shared address
space platform
(BBN Monarch, 1990)

SGI Origin 2000: NUMA Platform (1997)



J. Laudon and D. Lenoski. The SGI Origin: a ccNUMA highly scalable server.
Proc. of the 24th annual Intl. Symp. on Computer Architecture, Denver, 241 - 251, 1997

SGI Origin Network Topology

Example: “bristled hypercube” for 64 processors


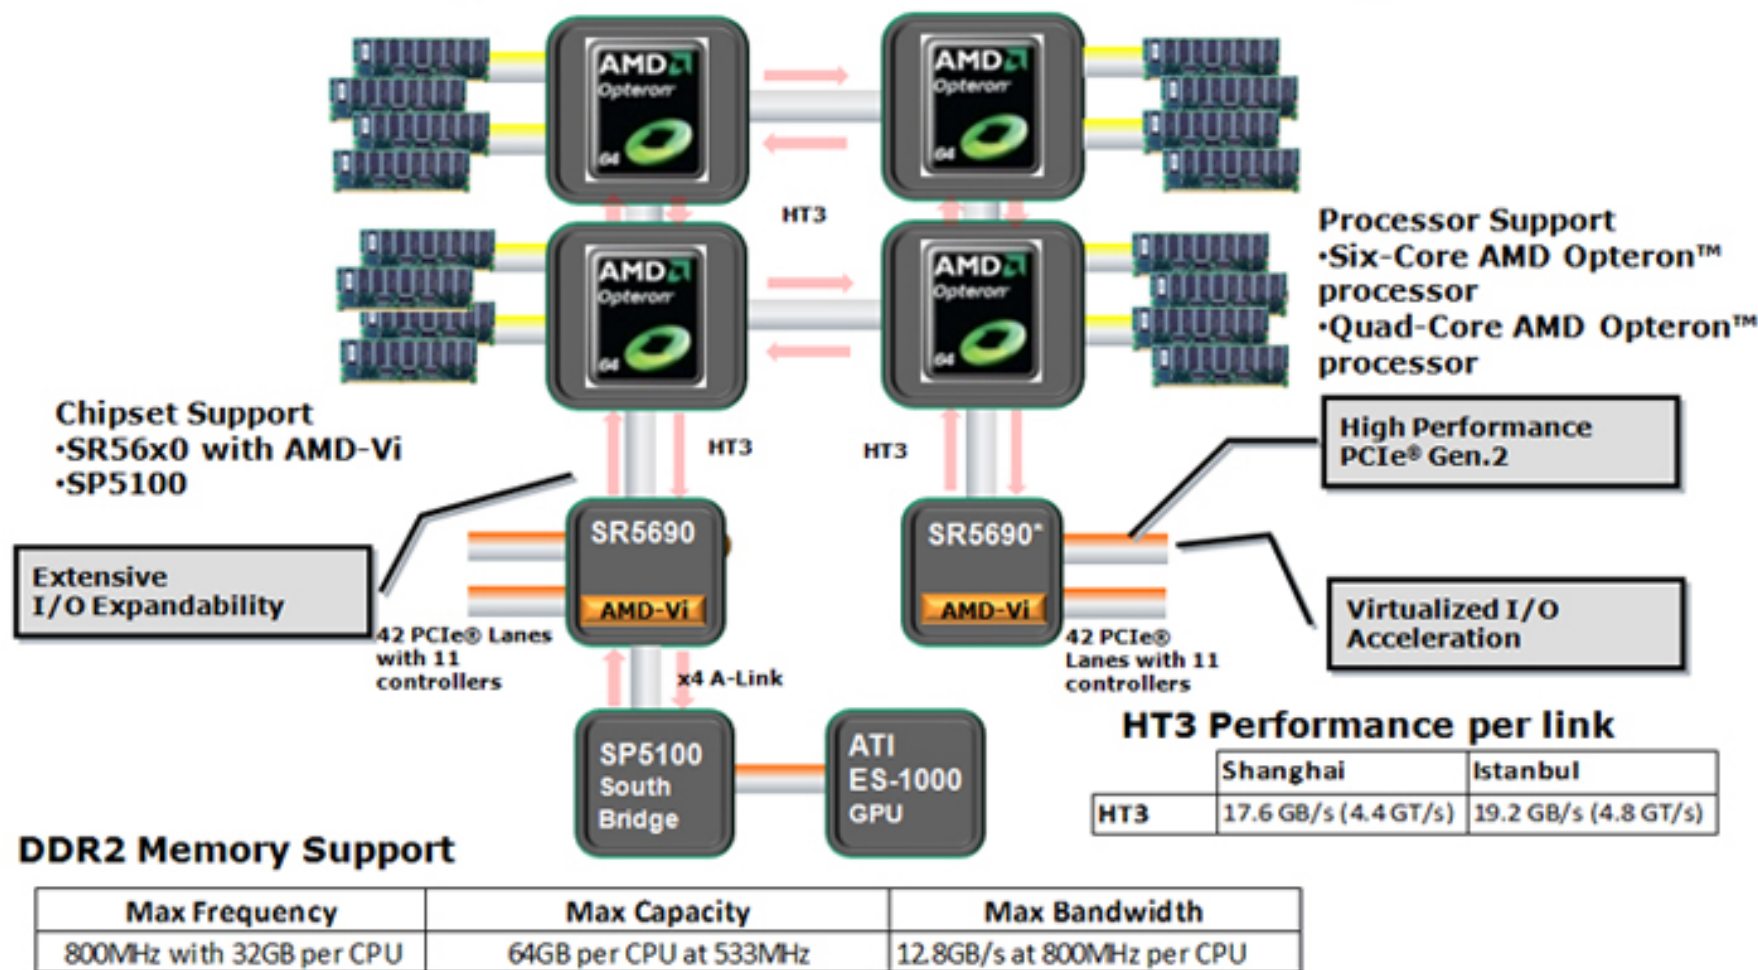


Figure credit: http://web.cecs.pdx.edu/~alaa/ece588/papers/laudon_isca_1997.pdf

Modern NUMA Multiprocessor

Six-Core AMD Opteron™ Processor with AMD Chipset Platform:



* Dual SR5690 is optional

Figure credit: AMD (<http://bit.ly/dT7Zxd>)

References

- Adapted from slides “Parallel Programming Platforms” by Ananth Grama accompanying course textbook
- Vivek Sarkar (Rice), COMP 422 slides from Spring 2008
- Jack Dongarra (U. Tenn.), CS 594 slides from Spring 2008, <http://www.cs.utk.edu/%7Edongarra/WEB-PAGES/cs594-2008.htm>
- Kathy Yelick (UC Berkeley), CS 267 slides from Spring 2007, http://www.eecs.berkeley.edu/~yelick/cs267_sp07/lectures
- Tom Henretty, Kevin Stock, Louis-Noël Pouchet, Franz Franchetti, J. Ramanujam and P. Sadayappan. Data Layout Transformation for Stencil Computations on Short-Vector SIMD Architectures. In ETAPS Intl. Conf. on Compiler Construction (CC'2011), Springer Verlag, Saarbrücken, Germany, March 2011.